

Recap: Representing numbers with float32



Key facts

8 bit biased exponent $\rightarrow \approx [10^{-38}, 10^{38}]$

24 bit fraction $\rightarrow$ **7 digit precision**

Recap: bfloat16



Key facts

8 bit biased exponent $\rightarrow \approx [10^{-38}, 10^{38}]$

7 bit fraction $\rightarrow$ 3 digit precision

Same as float32

Recap: int8




- Can represent $2^8 - 1 = 255$ values
- Range from $[-128, 127]$

Key facts

7 bits $\rightarrow [-128, 127]$

255 values only

Recap: Quantizing weights to int8



- **Quantization:** $\mathbf{W}_q = \text{round}(\gamma \cdot \mathbf{W})$ with $\gamma = \frac{2^{N-1}}{\max(|\mathbf{W}|)}$ and $N=8$
- **Dequantization:** $\hat{\mathbf{W}} = \frac{\mathbf{W}_q}{\gamma}$

This lecture


Explore various quantization techniques
and learn about their pros and cons

A closer look at model parameters

- Most computations within the model consist of

$$\mathbf{Y} = \mathbf{XW} + \mathbf{b}$$

- Let us ignore the bias (B) computation
- Consist of much smaller portion of the computation
 - $O(D^2)$ operations in $\mathbf{XW}$
 - $O(D)$ operations in $\mathbf{b}$
- Biases often not quantized




A closer look at model parameters


- Most computations within the model consist of

$$\mathbf{Y} = \mathbf{X} \mathbf{W}$$


- $\mathbf{X} \in \mathbb{R}^{B \times D_{in}}$ where B is the batch size
- Number of tokens could be another tensor dimensions
 - Will assume it to be 1 unless specified otherwise
- $\mathbf{W} \in \mathbb{R}^{D_{in} \times D_{out}}$
- The computation $\mathbf{XW}$ projects an input from D_{in} to D_{out} dimensions




A recap of matrix multiplication




A recap of matrix multiplication




A recap of matrix multiplication




A recap of matrix multiplication



A recap of matrix multiplication




A recap of matrix multiplication



Repeat for the next prompt in the batch


A recap of matrix multiplication




Repeat for the next prompt in the batch

A recap of matrix multiplication


$$\mathbf{X} \in \mathbb{R}^{B \times D_{in}}$$



$$\mathbf{W} \in \mathbb{R}^{D_{in} \times D_{out}}$$




$$\mathbf{Y} \in \mathbb{R}^{B \times D_{out}}$$



Back to quantization

Quantize weights only



Key idea


Only one, or few, weights
dequantized at a time

Everything else remains in int8

Also quantize activations

bf16

$$\mathbf{X} \in \mathbb{R}^{B \times D_{in}}$$




quantize



int8


$$\hat{\mathbf{X}} \in \mathbb{R}^{B \times D_{in}}$$



x

int8


$$\hat{\mathbf{W}} \in \mathbb{R}^{D_{in} \times D_{out}}$$



=

int8

$$\hat{\mathbf{Y}} \in \mathbb{R}^{B \times D_{out}}$$



Note

- $\hat{\cdot}$ in $\hat{\mathbf{X}}$, $\hat{\mathbf{W}}$ and $\hat{\mathbf{Y}}$ denotes the quantized version
- $\mathbf{X}$ and $\mathbf{Y}$ are both activations
- Input of one **layer** L , that is, $\mathbf{X}$, is output of **layer** $L-1$

Key idea

All operations in int8

GPU can leverage custom integer kernels

Weight + Activation quantization

- Activation quantization adds more complications

Weight + Activation quantization

- Activation quantization adds more complications
- Weights are fixed, that is, they do not change once we are done training
 - Quantization **scale s** needs to be learnt **only once**
- Activations change with every new prompt
 - Need to set the quantize scale s for **every new input**

$$\gamma = \frac{2^{N-1}}{\max(|\mathbf{X}|)}$$

Change for every prompt

$$\gamma = \frac{2^{N-1}}{\max(|\mathbf{W}|)}$$

Fixed per model

Weight + Activation quantization

- Activation quantization adds more complications
- Weights are fixed, that is, they do not change once we are done training
 - Quantization **scale s** needs to be learnt **only once**
- Activations change with every new prompt
 - Need to set the quantize scale s for **every new input**
- **Dynamic quantization**
 - Learn the scale for every input on the fly
 - Accurate but slow (have to re-compute scale for every input)

$$\gamma = \frac{2^{N-1}}{\max(|\mathbf{X}|)}$$

Change for every prompt

$$\gamma = \frac{2^{N-1}}{\max(|\mathbf{W}|)}$$

Fixed per model

Weight + Activation quantization

- Activation quantization adds more complications
- Weights are fixed, that is, they do not change once we are done training
 - Quantization **scale s** needs to be learnt **only once**
- Activations change with every new prompt
 - Need to set the quantize scale s for **every new input**
- **Dynamic quantization**
 - Learn the scale for every input on the fly
 - Accurate but slow (have to re-compute scale for every input)
- **Static quantization**
 - Use a calibration dataset (e.g., few hundred samples) to learn the scale
 - Fast but less accurate

$$\gamma = \frac{2^{N-1}}{\max(|\mathbf{X}|)}$$


Change for every prompt

$$\gamma = \frac{2^{N-1}}{\max(|\mathbf{W}|)}$$

Fixed per model


LLM.int8() quantization

- Available in HuggingFace via the [bitsandbytes](#) library
- **Outliers** emerge after model size exceeds ~3B
- Luckily, outliers concentrated in **few channels**
- **Key idea:** Mixed precision quantization
 - Do not quantize outliers (less than 1% of the weights)
 - All other weights get quantized to int8




Mixed precision quantization


$$\mathbf{X} \in \mathbb{R}^{D_{in}}$$



$$\mathbf{W} \in \mathbb{R}^{D_{in} \times D_{out}}$$



$$\mathbf{Y} \in \mathbb{R}^{D_{out}}$$






Mixed precision quantization


$\mathbf{X} \in \mathbb{R}^{D_{in}}$

$\mathbf{W} \in \mathbb{R}^{D_{in} \times D_{out}}$

$\mathbf{Y} \in \mathbb{R}^{D_{out}}$




$\mathbf{X}$




$\mathbf{x}$

$\mathbf{W}_{[:,0]}$




$=$

Σ




$=$




$\mathbf{Y}_{[0,0]}$

$\mathbf{X}$





$\mathbf{x}$

$\mathbf{W}_{[:,0]}$




$=$


Σ




$\mathbf{x}$



$=$



$=$



$\mathbf{Y}_{[0,0]}$

Re-
arrange
No effect
on output

Quantize
to int8

Keep
in bf16

Outlier activations and
corresponding weights.
E.g., all activations are
max. 1.1 but outliers are
>10

LLM.int8(): Pros and cons


- Simple to implement
- Works on the fly
- Need to set the outlier threshold manually (e.g., 6)
- Paper shows that the runtime is slower than the base model
- May not work on all hardware — bitsandbytes works on GPUs only (so far)

GPTQ

- LLM.int8() performs poorly when scaled to 4 bits
- GPTQ aims to quantize to lower than 8-bit precision, e.g., 4 or 2 bits
- **Key idea:** Minimize the loss between the output of quantized and unquantized matmul



$$\hat{\mathbf{W}} = \arg \min || \mathbf{XW} - \mathbf{X}\hat{\mathbf{W}} ||_2^2$$

GPTQ: Performing quantization in parallel




Can perform the computation for different columns **in parallel**

$$\begin{aligned}
 \hat{\mathbf{W}} &= \arg \min ||\mathbf{XW} - \mathbf{X}\hat{\mathbf{W}}||_2^2 \\
 &= \arg \min \sum ||\mathbf{XW}_{[:,i]} - \mathbf{X}\hat{\mathbf{W}}_{[:,i]}||_2^2
 \end{aligned}$$



Iteratively quantizing weights



- Quantize this weight
- Update the unquantized weights in $\mathbf{W}_{[:,0]}$ to minimize $||\mathbf{XW} - \mathbf{X}\hat{\mathbf{W}}||_2^2$
- Repeat for all weights in $\mathbf{W}_{[:,0]}$
- Recursively repeat for all columns



GPTQ: Pros and cons

- Quantizes models very quickly by using second order optimization
- Only quantizes weights so runs on the fly
- Needs a calibration dataset
- Unlike LLM.int8(), does not need to guess the outlier values, instead, works based off the data

Exercise

SmoothQuant

- **Intuition:** No outliers in weights, many large outliers in activations
 - **Weights:** No outliers means we use the quantization range **efficiently** (many effective bits)
 - **Activations:** Large outliers means we use the quantization range **inefficiently** (few effective bits)
- However, outliers are concentrated in a few channels



SmoothQuant

- **Key idea:** Migrating outliers between activations and weights
 - Smooth out the activations by dividing them with a channel-wise scaling term S
 - Divide the weights by the same term to migrate the outliers


- $\bar{\mathbf{X}} = \mathbf{S}^{-1}\mathbf{X}$ with $\mathbf{X}, \mathbf{S}, \bar{\mathbf{X}} \in \mathbb{R}^{D_{in}}$ and $S > 1$

- Obtain $\bar{\mathbf{W}}$ by multiplying with $\mathbf{S}$


- $\mathbf{XW} = \bar{\mathbf{X}}\bar{\mathbf{W}}$

- $S_j = \max(|\mathbf{X}_j|)^\alpha / \max(|\mathbf{W}_j|)^{1-\alpha}$
 - Pick $\alpha \in [0,1]$ based on the model / quantization


- Proceed with quantization



(a) Original



SmoothQuant



SmoothQuant: Pros and cons

- The smoothing process done offline so faster than dynamic quantization
- Needs a calibration dataset
- Can mix and match with downstream quantization strategies


AWQ

- Somewhat of a mix of ideas discussed earlier
- **Key ideas**
 - Weights and activations interact, so quantization should be mindful of that
 - Scale weights and activations before quantization
 - Learn the quantization parameters by solving an optimization problem over the data
 - Avoid mixed precision like mix of int8 and bf16

$$\mathbf{s}^* = \arg \min_{\mathbf{s}} \mathcal{L}(\mathbf{s})$$

$$\mathcal{L}(\mathbf{s}) = \|Q(\mathbf{W} \cdot \text{diag}(\mathbf{s}))(\text{diag}(\mathbf{s})^{-1} \cdot \mathbf{X}) - \mathbf{W}\mathbf{X}\|$$




Quantization function
e.g., int4



Scaling factor
(per channel)



AWQ



AWQ: Pros and cons


- Needs a calibration dataset
- Weight only quantization (doesn't quantize activations)
- No mixed precision so avoids limitations of LLM.int8()
- Works well is very low precision, e.g., int4

Key considerations in quantization

- Do we need a **calibration dataset**?
- Is this a **dynamic** or **static** quantization strategy?
- Are only **weights** getting quantized or **activations** too?
 - Having both in int8 can leverage integer kernels to provide speedups
- Can we **learn from data** or do we need to make **heuristic based choices**?
- Does my strategy have **hardware support**?
 - Even int8 is not supported on all hardware
 - Some hardware does not support bf16
- **Space savings** and **computational speed ups** are two different things
 - Not all quantization methods will lead to a speedup, some will slow you down, e.g., LLM.int8()


Looking beyond quantization

Magnitude pruning




- **Key idea:** Prune weights that are lower than a certain threshold

WandA (weights and activations) pruning



- **Key idea:** Prune weights that have a low importance based on the $|\mathbf{W}| \cdot \|\mathbf{X}\|_2$ metric
- Similar idea to AWQ and LLM.int8(): Weight importance is determined by the activations too

Structured pruning



- **Key idea:** Fixed ratio of weights (e.g., 2:4) is pruned between each structure
- Can leverage hardware acceleration

Exercise

References

- [Making LLMs lighter with AutoGPTQ and transformers](#)
- [Introduction to Weight Quantization](#)
- [A Gentle Introduction to 8-bit Matrix Multiplication for transformers at scale using Hugging Face Transformers, Accelerate and bitsandbytes](#)
- [A Visual Guide to Quantization](#)
- [SmoothQuant: Accurate and Efficient Post-Training Quantization for Large Language Models](#)
- [GPTQ: Accurate Post-Training Quantization for Generative Pre-trained Transformers](#)
- [AWQ: Activation-aware Weight Quantization for LLM Compression and Acceleration](#)
- [LLM.int8\(\): 8-bit Matrix Multiplication for Transformers at Scale](#)
- [A Simple and Effective Pruning Approach for Large Language Models](#)
- [Accelerating Sparse Deep Neural Networks](#)
- [MIT 6.5940: TinyML and Efficient Deep Learning Computing](#)