

# Scaling LLM Training: Part I

Vedant Nanda  
Researcher @ Aleph Alpha Research



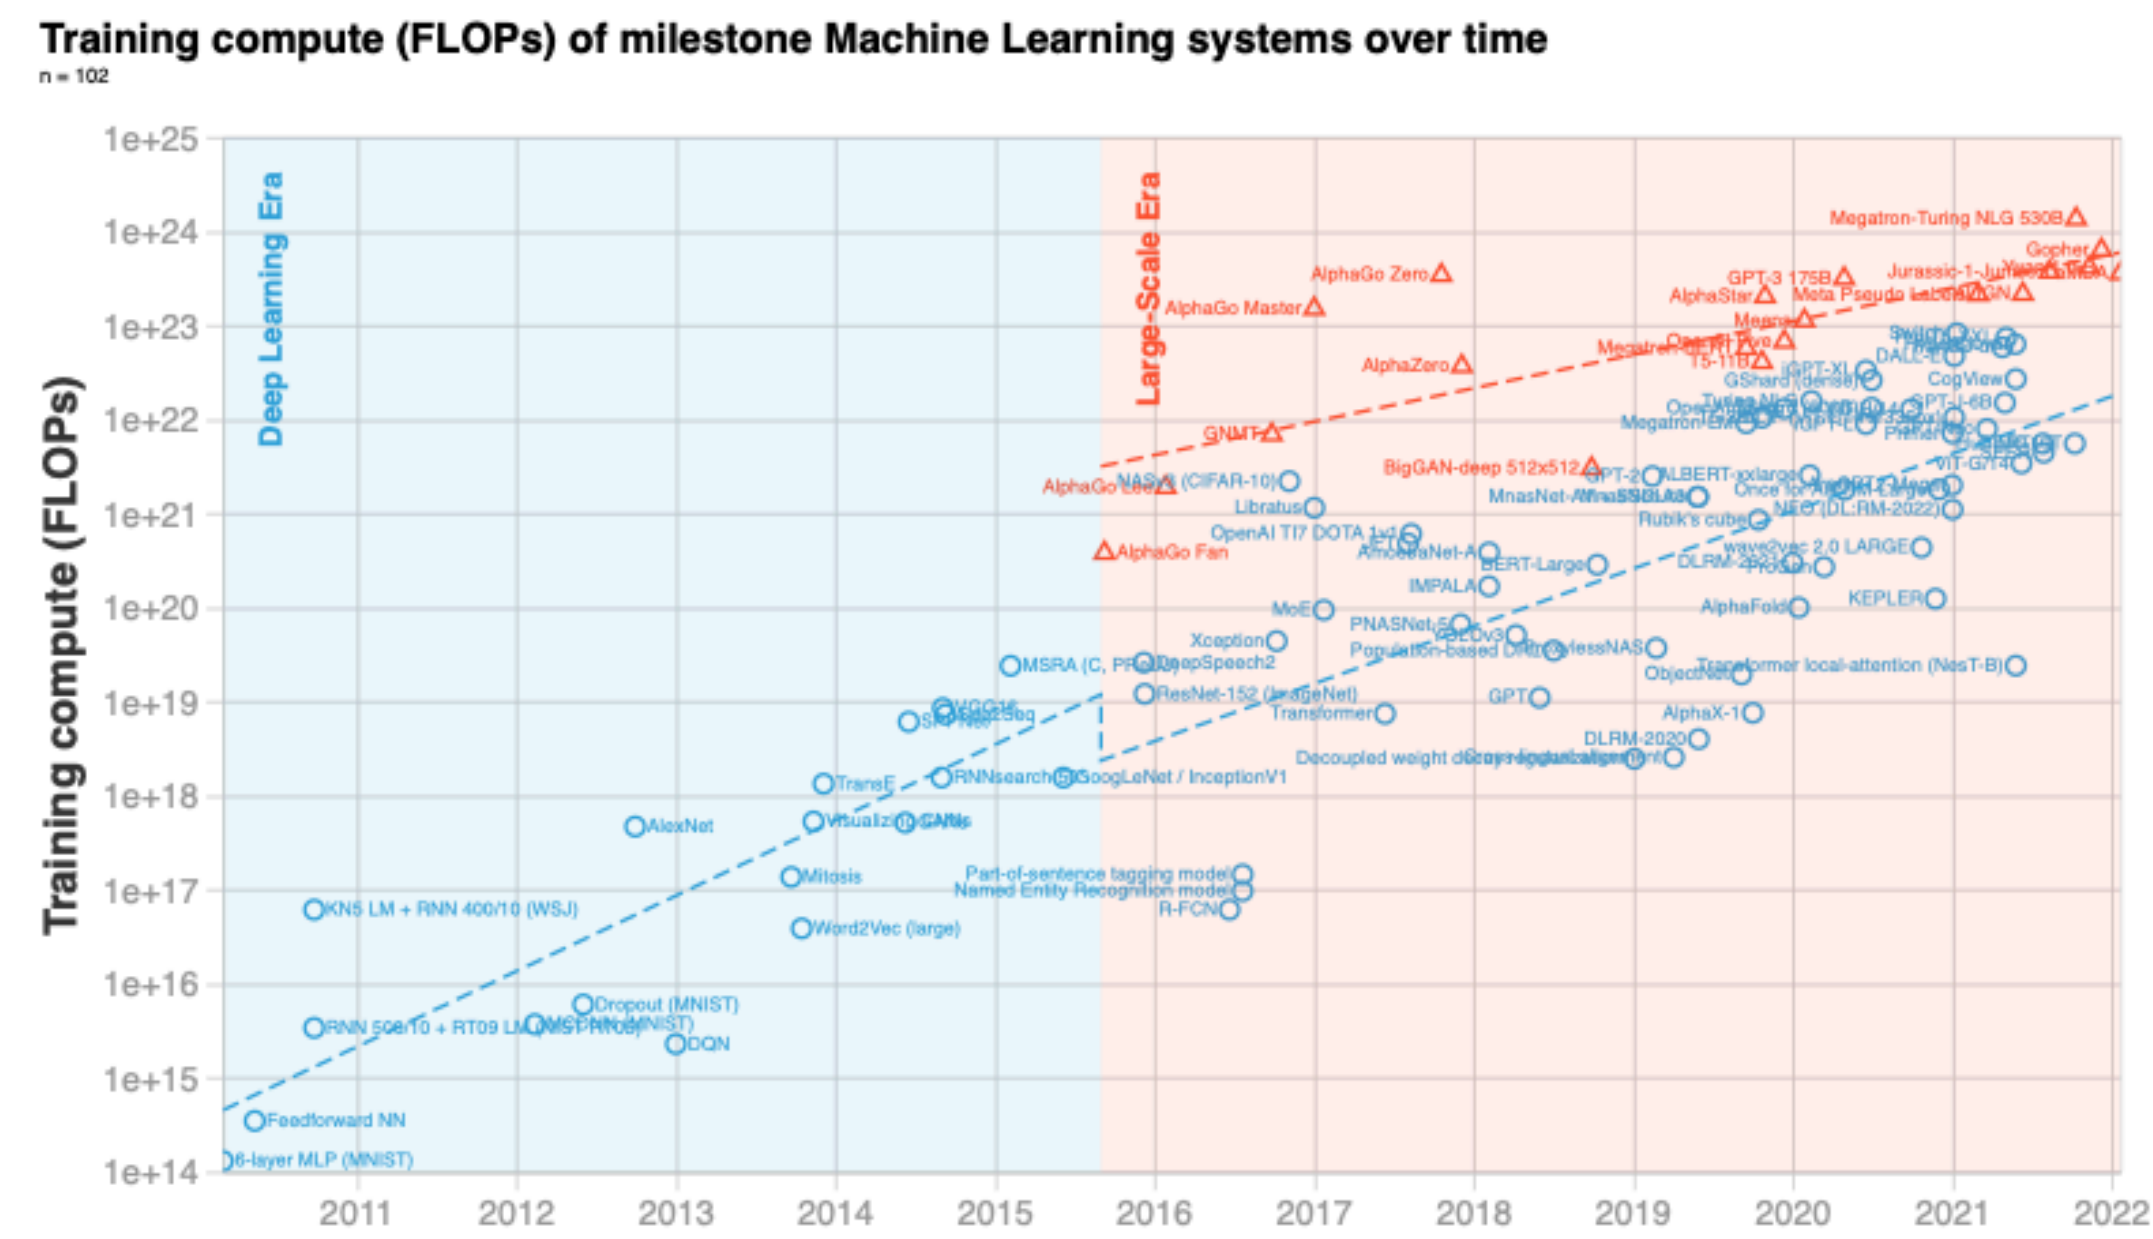
# About Me

- PhD in CS from University of Maryland and MPI-SWS
- At Aleph Alpha for the last year: optimizing inference and pre-training
- We released base and instruct tokenizer-free models at ICLR last month
  - <https://huggingface.co/collections/Aleph-Alpha>
- Alt title for this lecture: stuff I wish I knew about pre-training LLMs

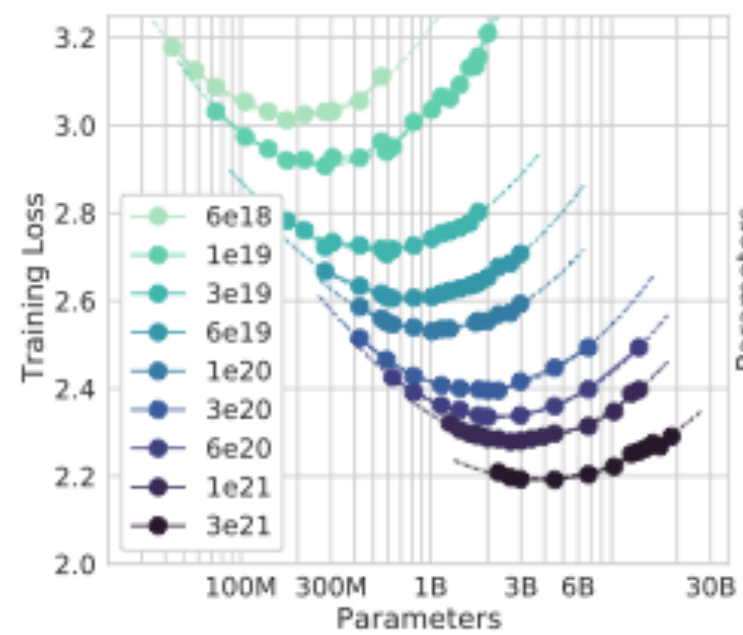
# This Lecture

- Motivation (5 min)
- Deep dive into transformer compute (20 min)
- Understanding training workloads (20 min)
- Hands-on exercise (45 min)
- Scaling LLMs on one GPU: activation checkpointing (10 min)
- Scaling LLMs on one GPU: gradient accumulation (10 min)
- Scaling beyond one GPU: data parallelism (20 min)
- Hands-on exercise (45 min)

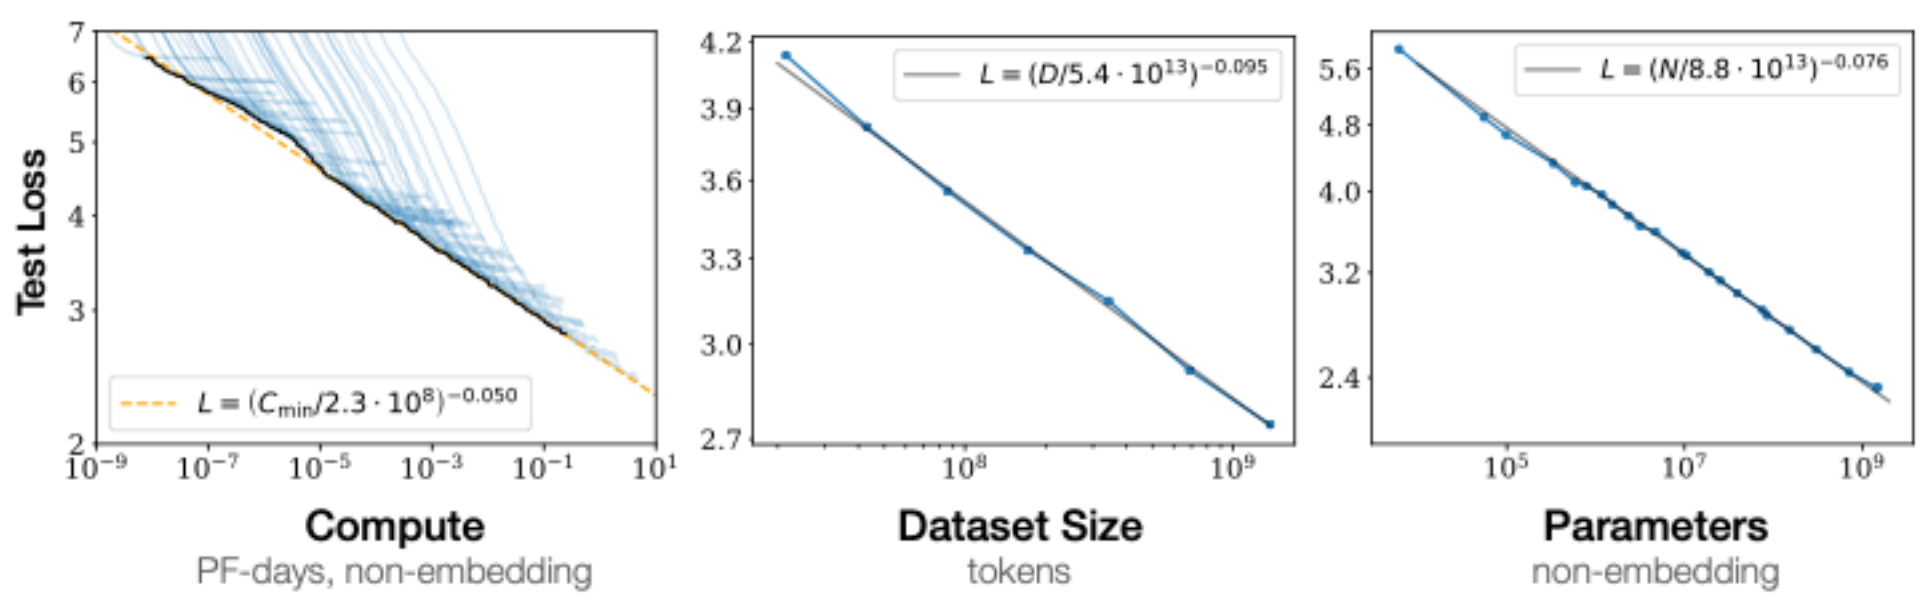
# Scale: A Key Ingredient in SOTA LLMs



“Compute Trends Across Three Eras of Machine Learning” Sevilla et al. 2022



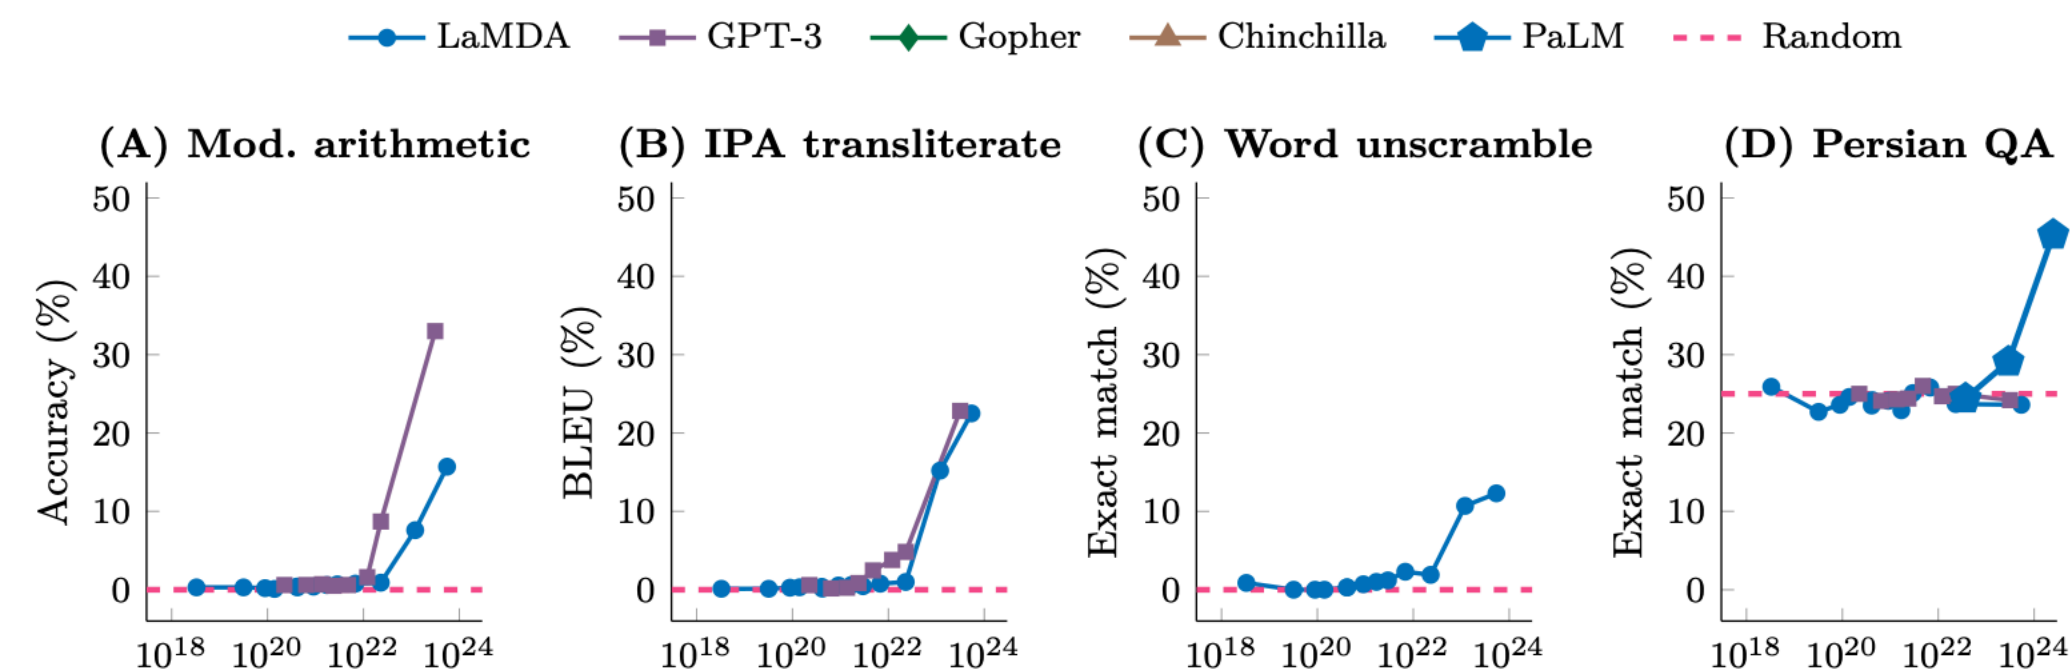
“Training Compute-Optimal Large Language Models” Hoffmann et al. 2022



“Scaling Laws for Neural Language Models” Kaplan et al. 2022

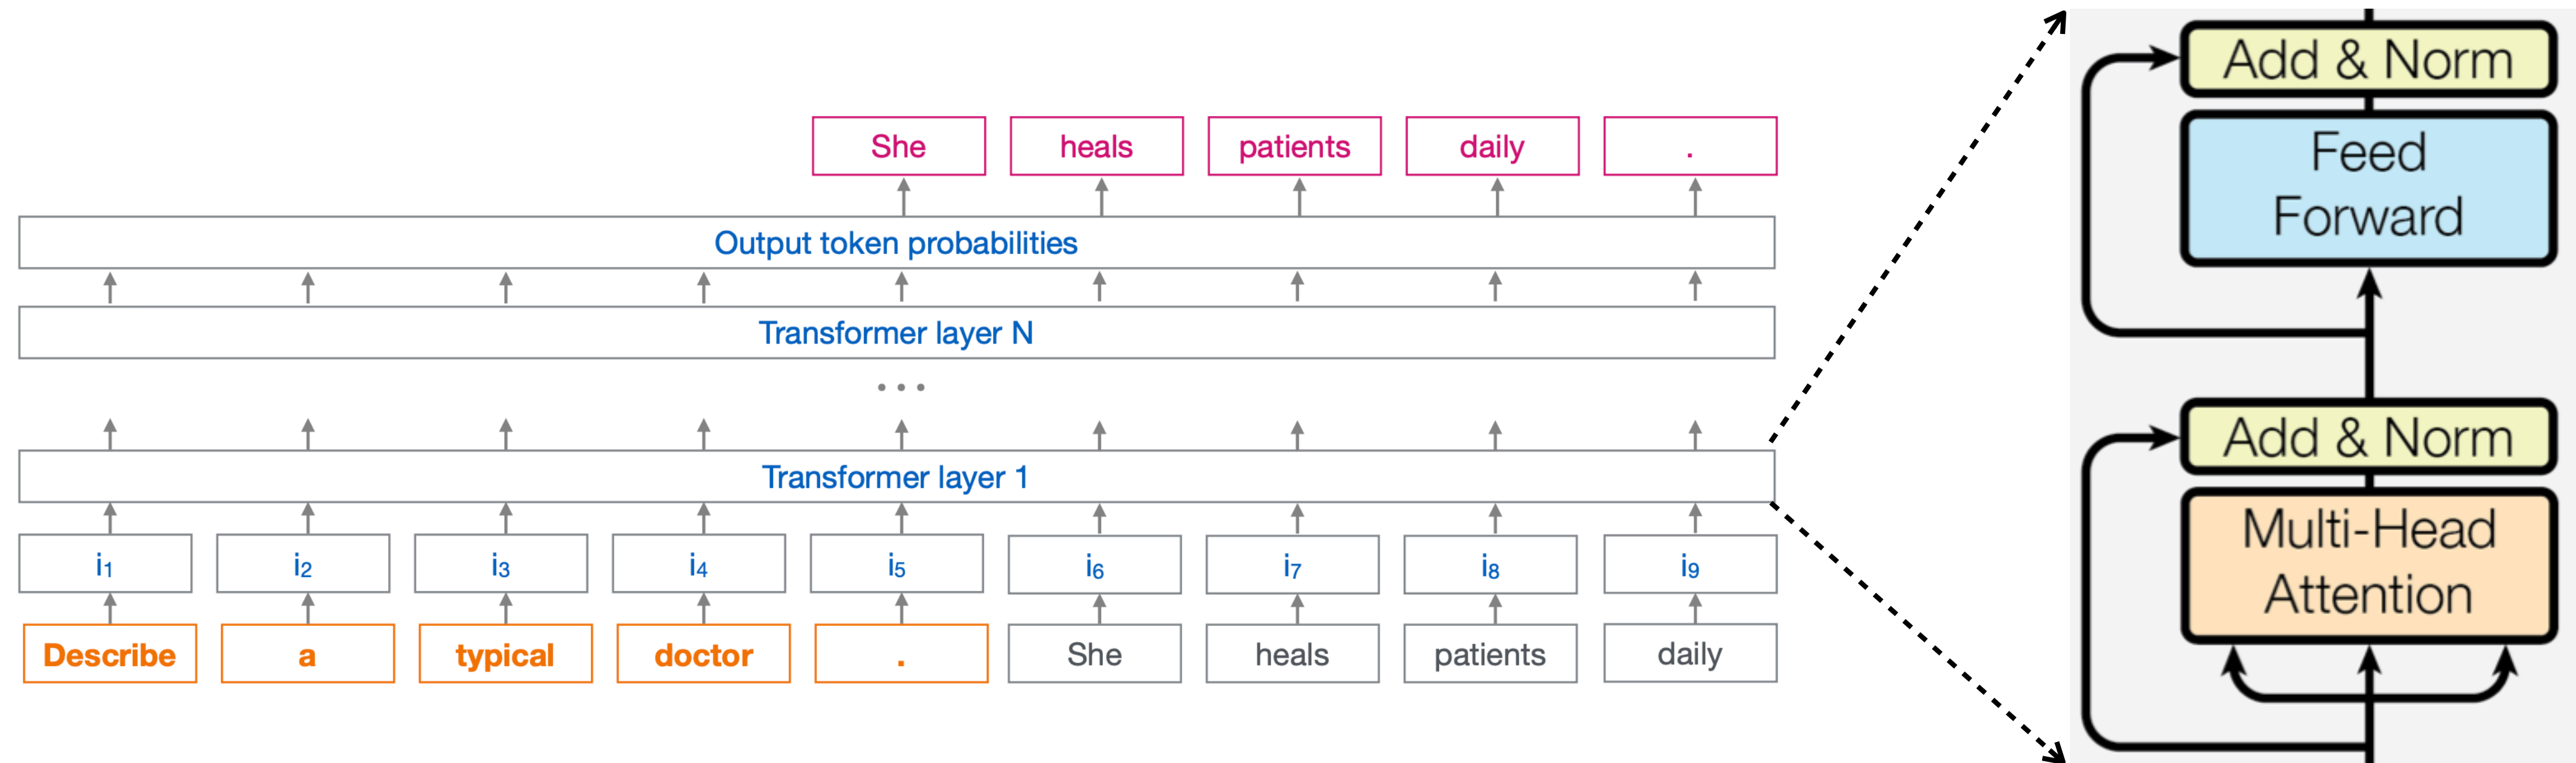
# Why Should ML Researchers Care?

- Arguably the most pivotal Deep Learning paper (AlexNet circa 2012) was a systems paper that trained a CNN with model parallelism
- Many architectures died because they didn't scale (eg: LSTM)
  - *The biggest lesson that can be read from 70 years of AI research is that general methods that leverage computation are ultimately the most effective, and by a large margin.* — Bitter Lesson by Rich Sutton
- Emergence at scale; many innovations today are at the edge of hardware

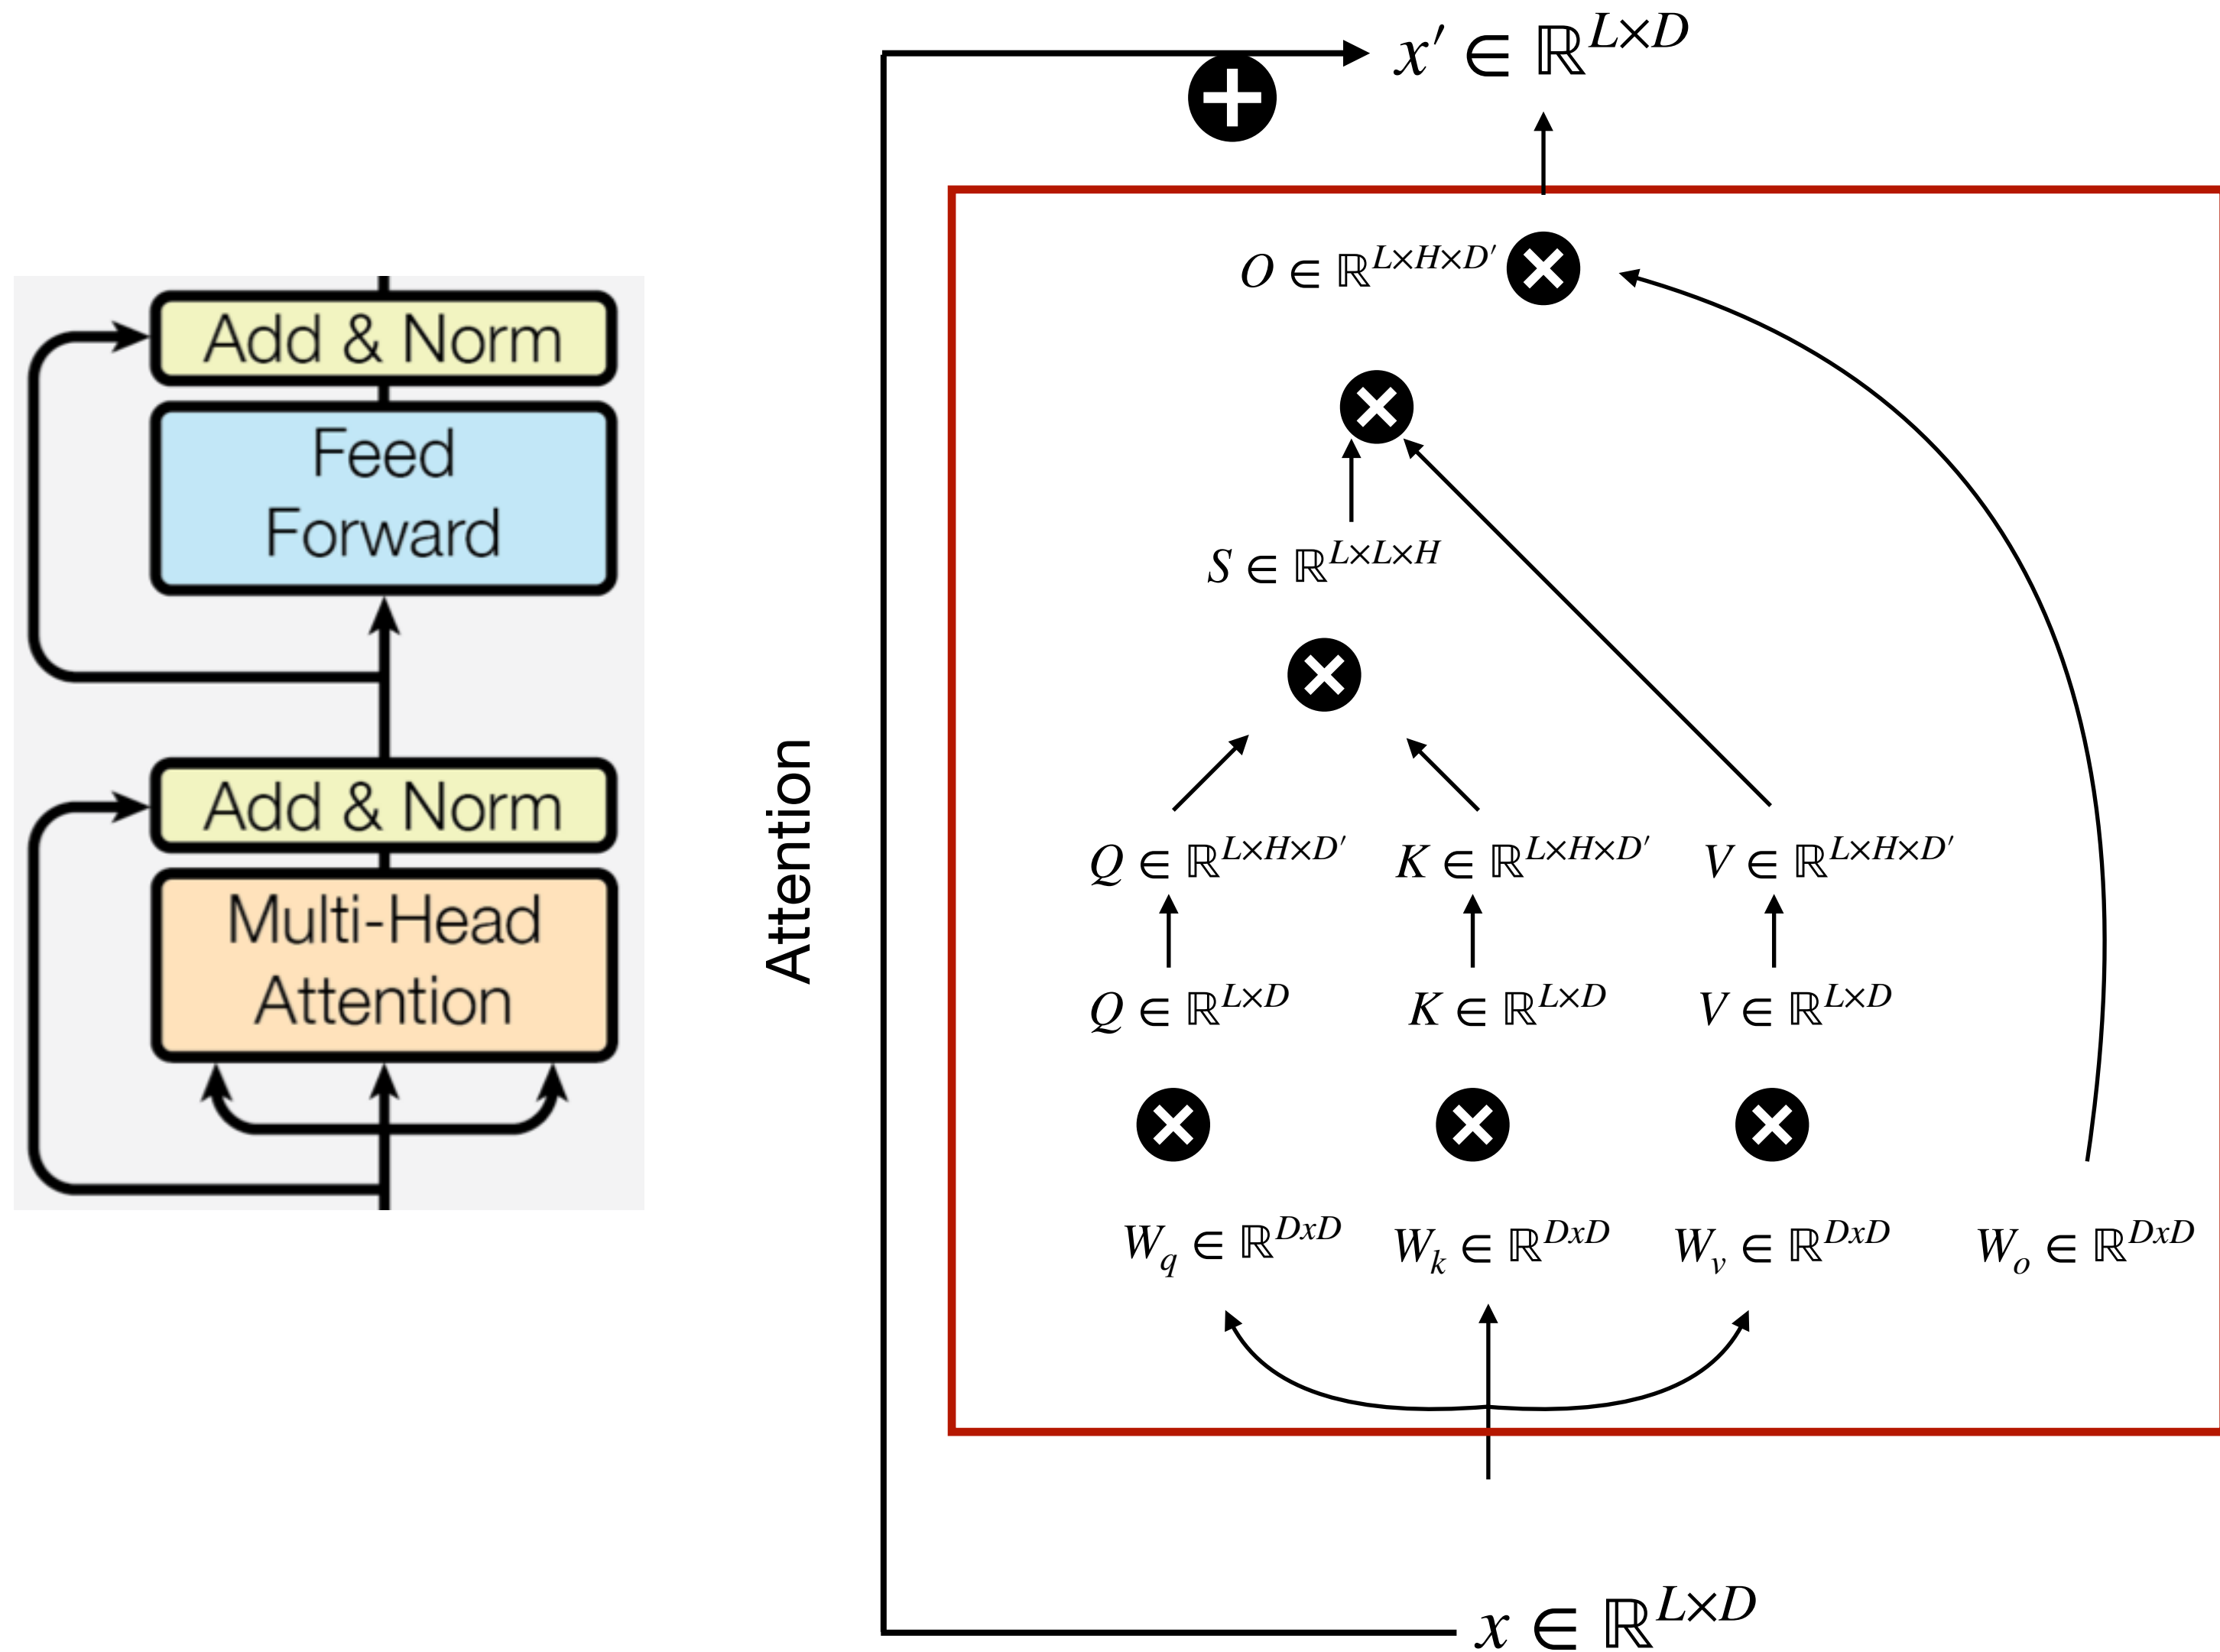


Emergent Abilities of Large Language Models, 2022

# Transformers Recap

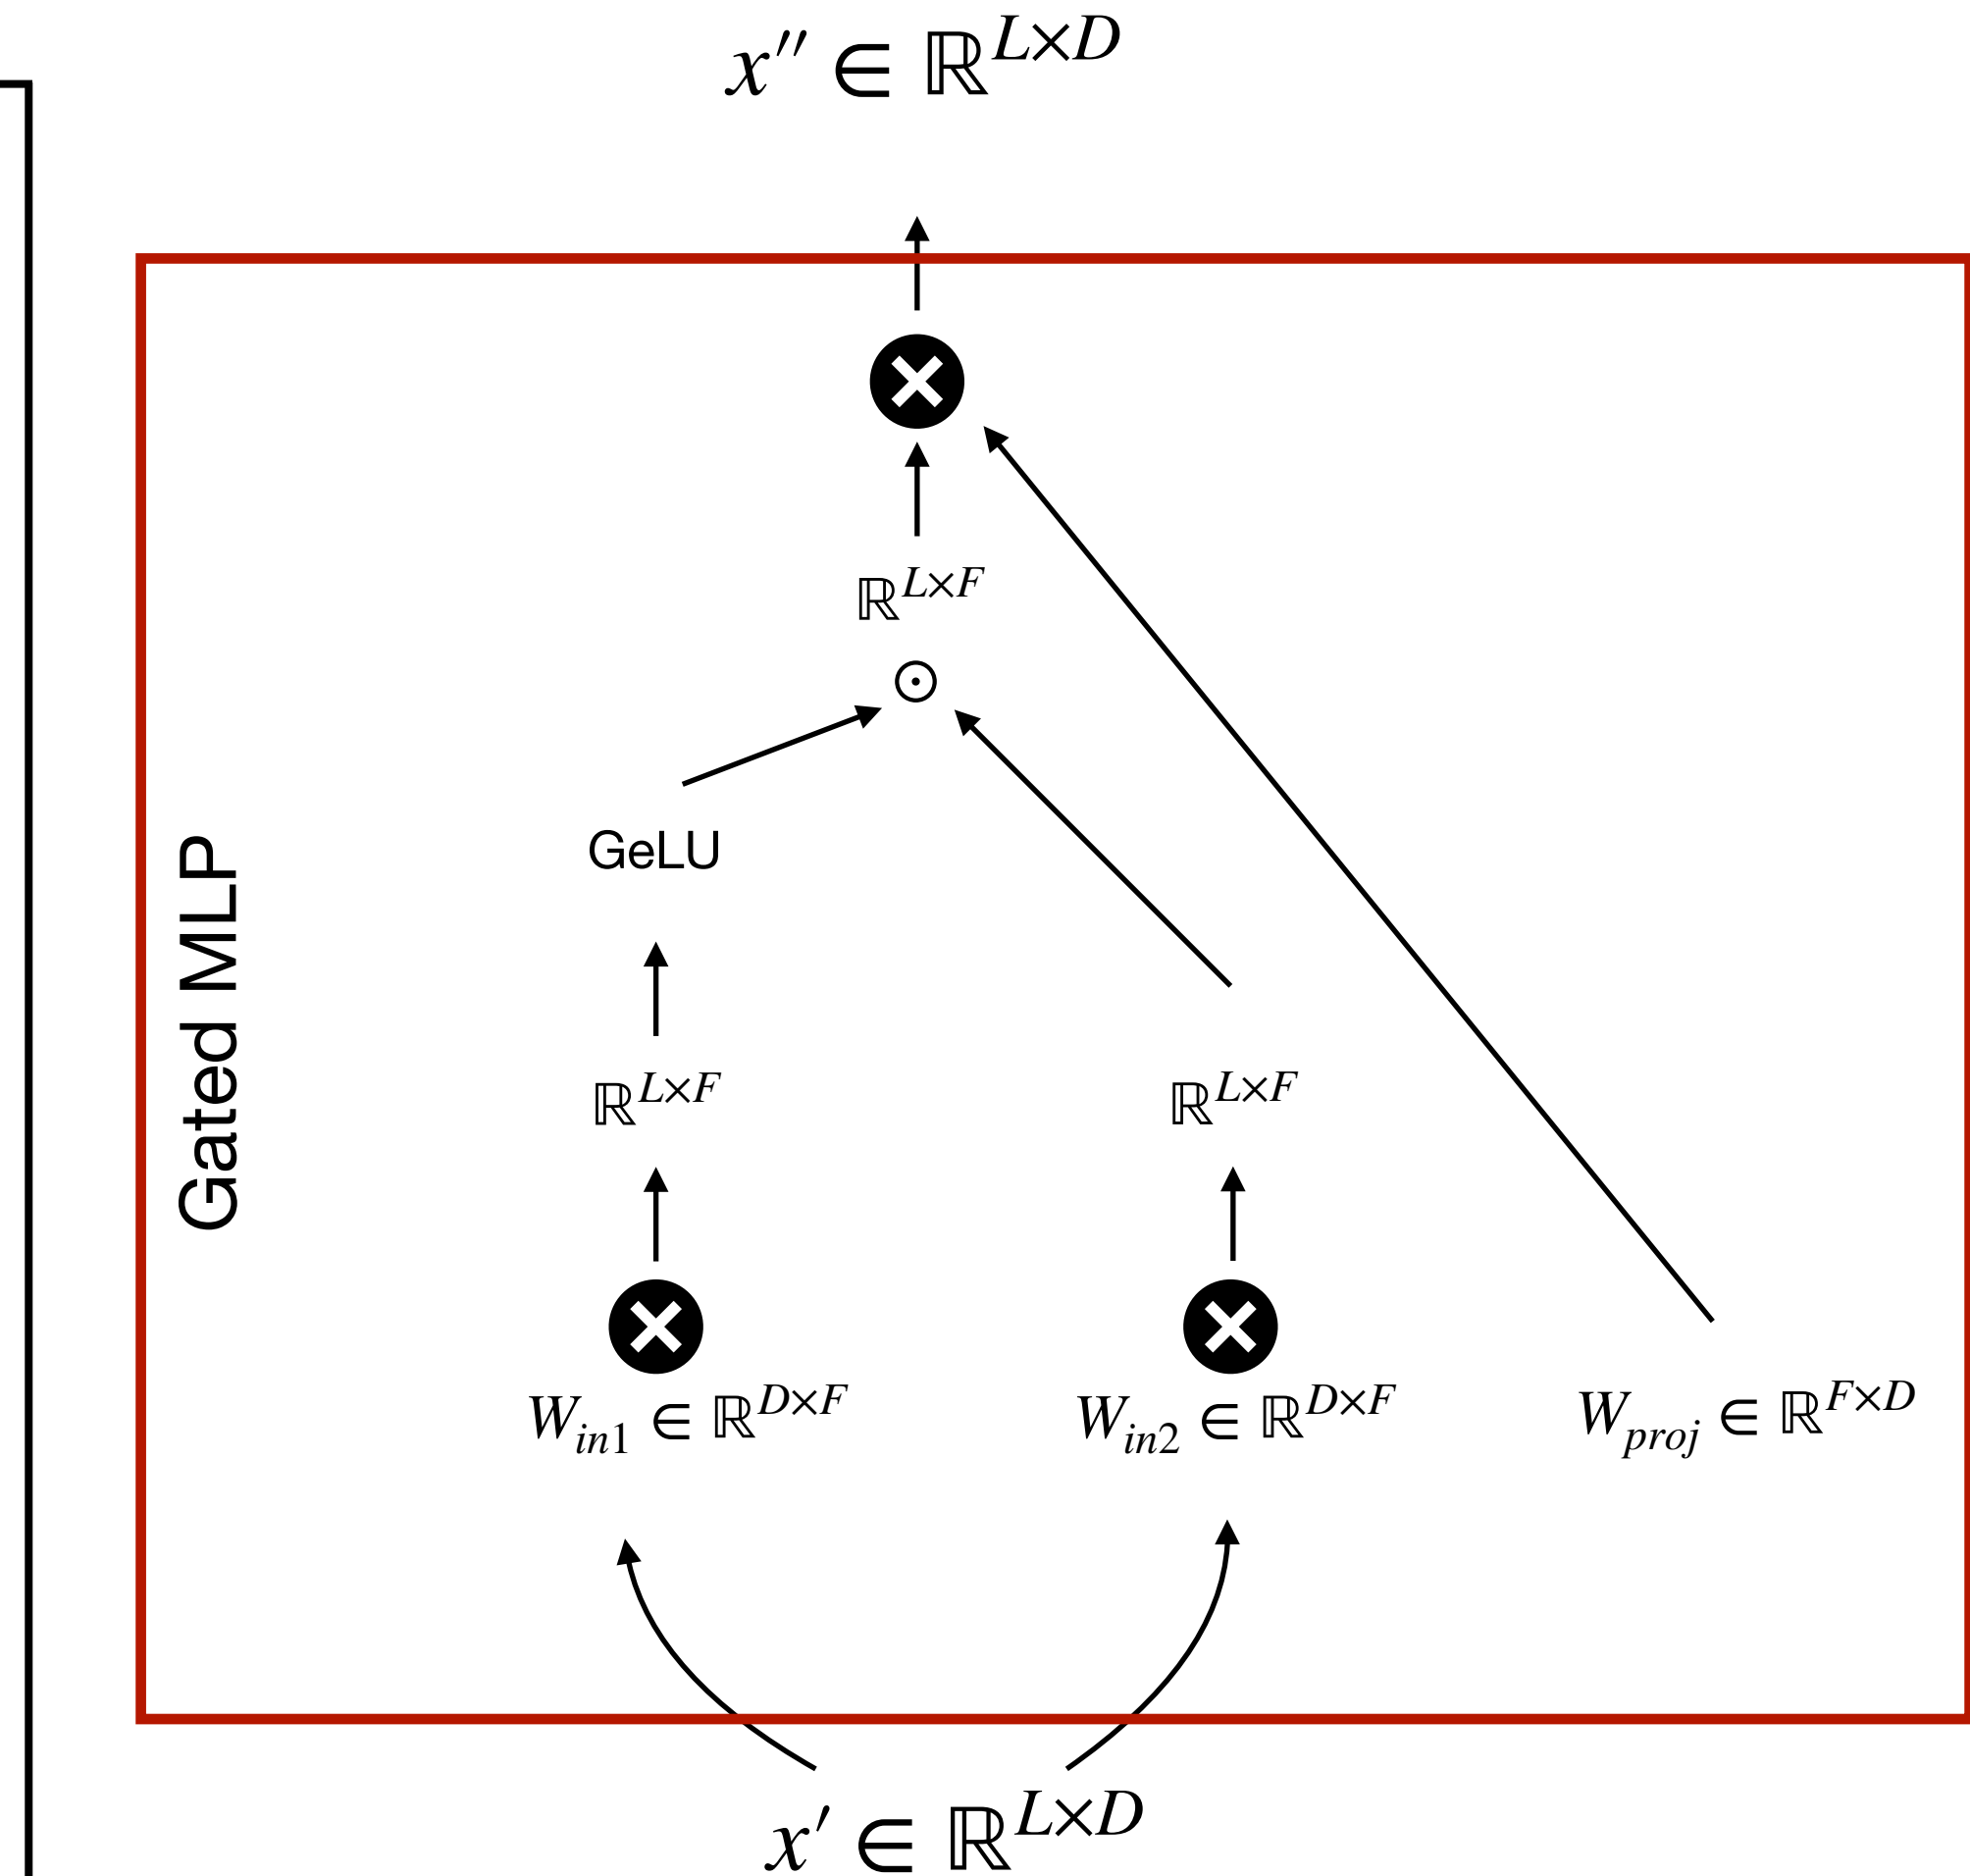
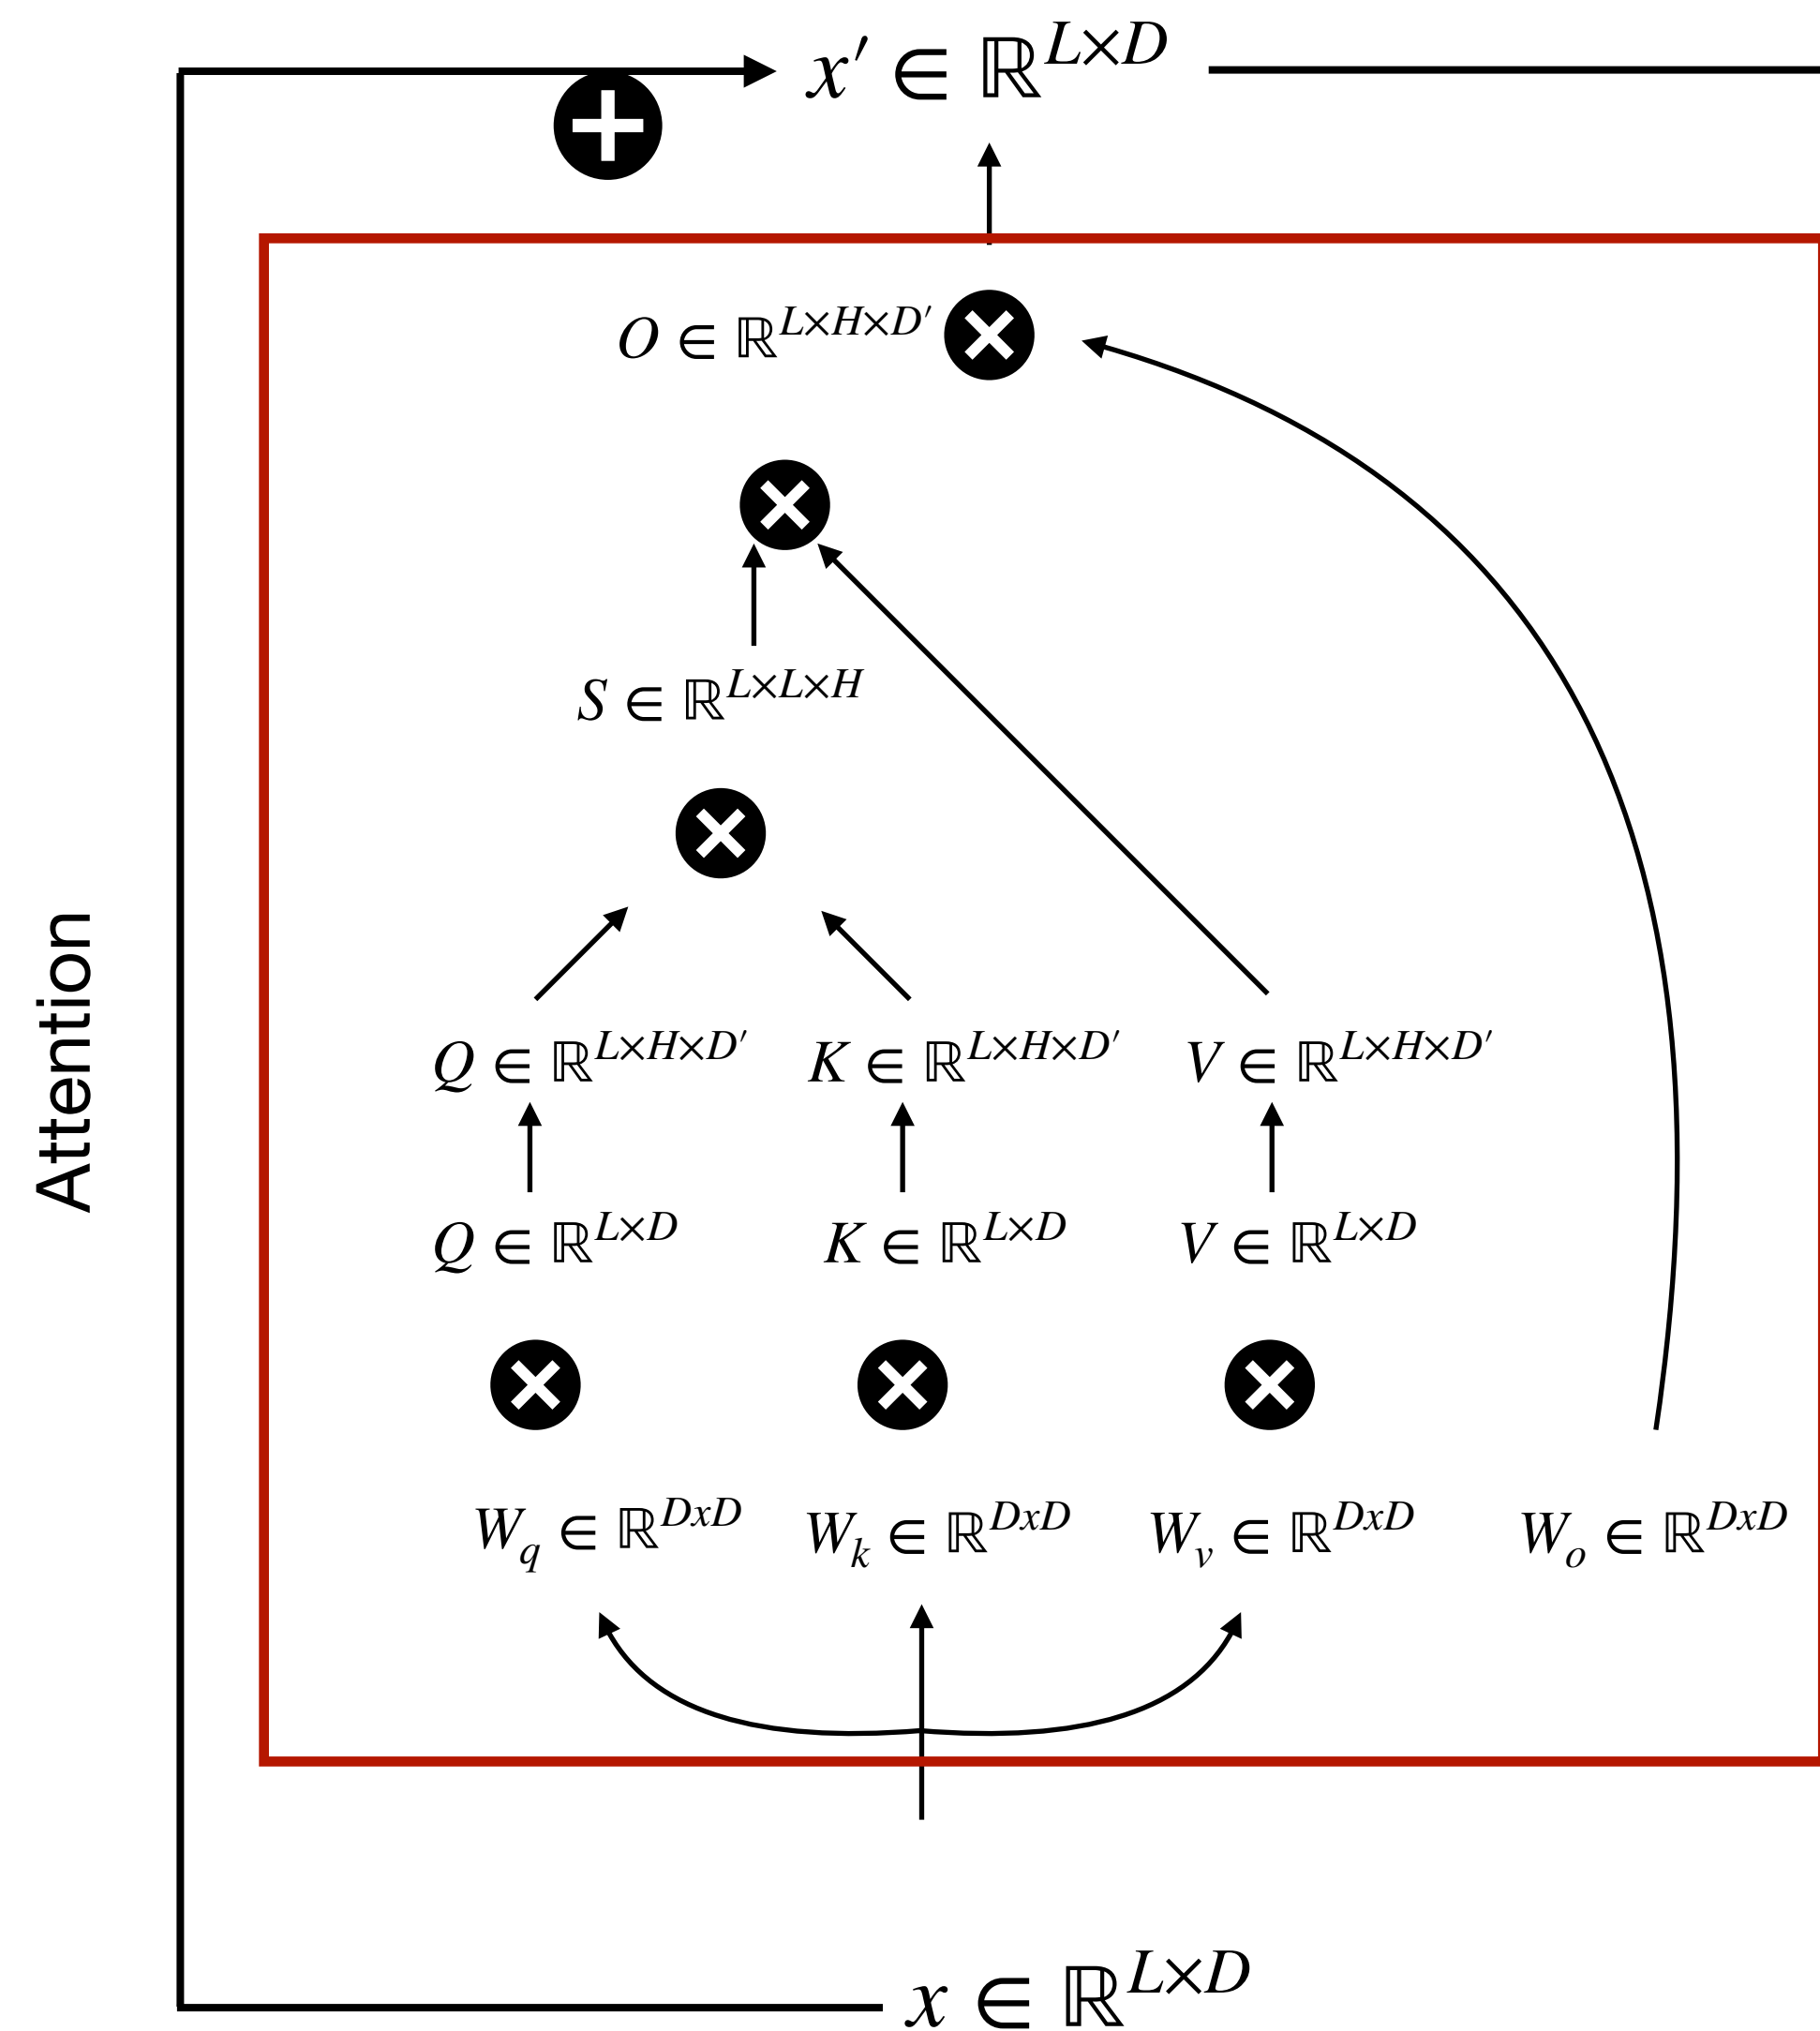


# Transformers: What Do We Compute?





# Transformers: What Do We Compute?





# Transformers: What Do We Compute?

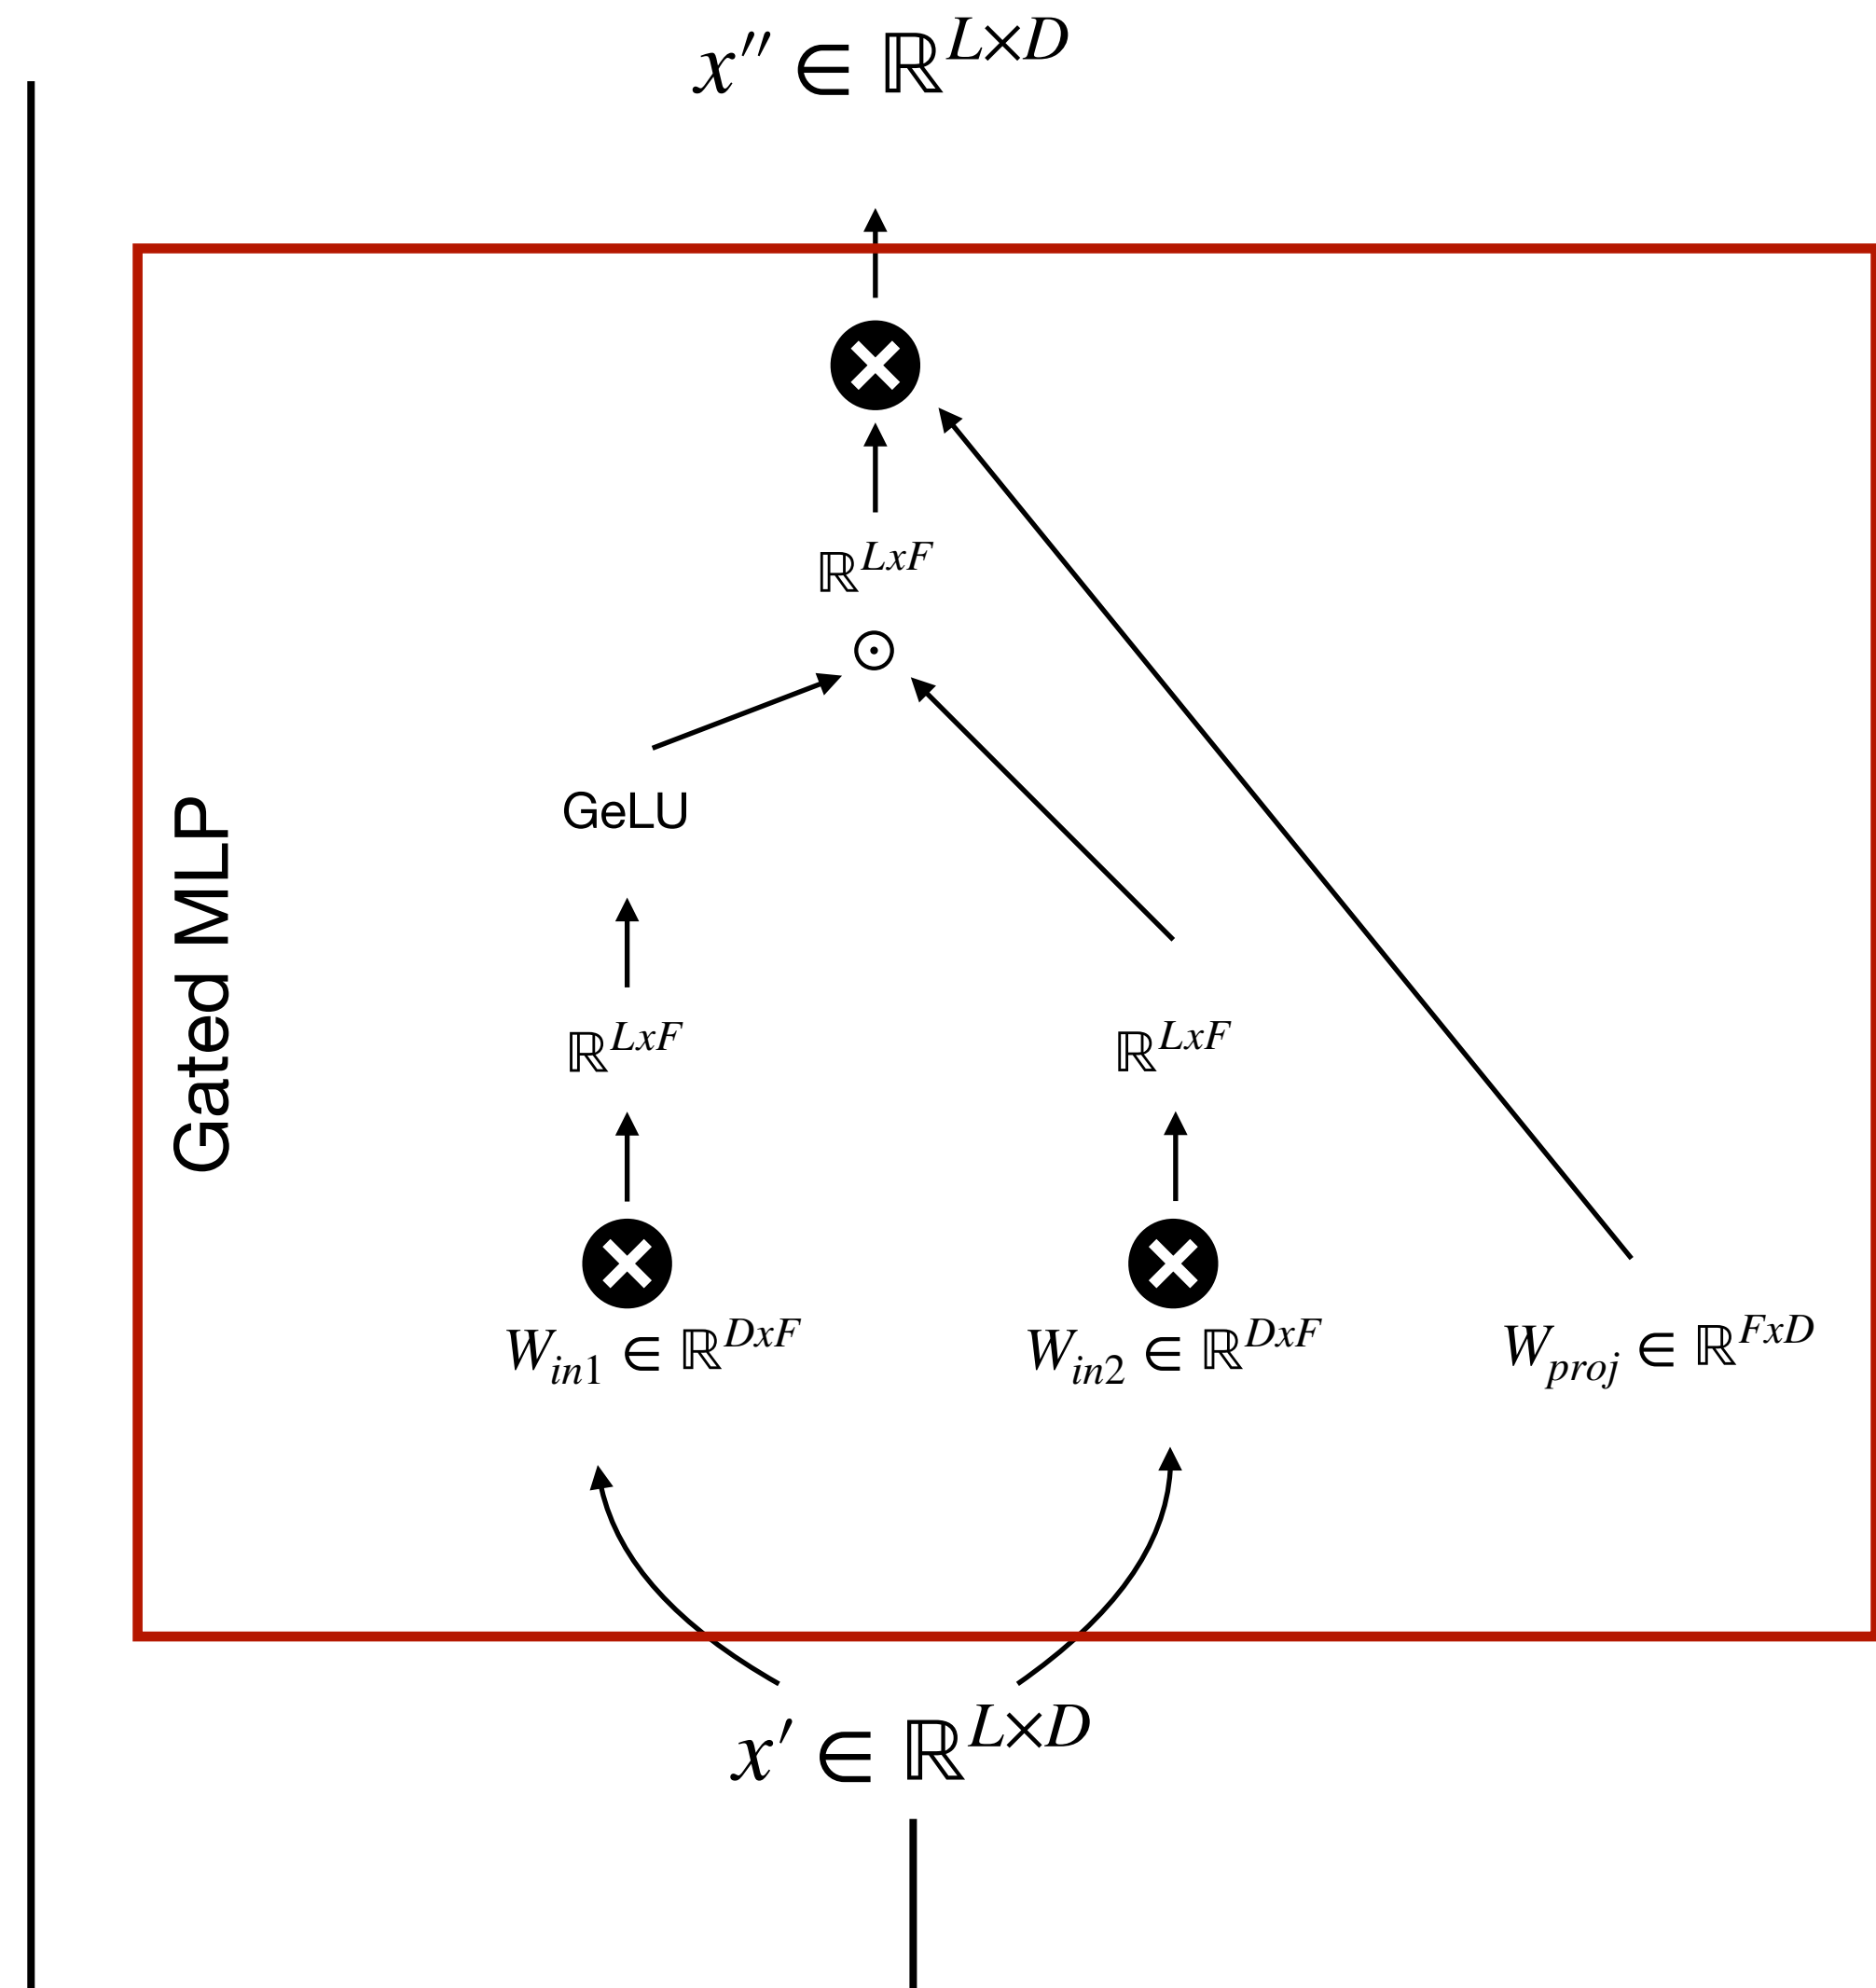
## Why gating?

Standard choice in all latest architectures:  
Llama, DeepSeek, Qwen, Mistral

### 4 Conclusions

We have extended the GLU family of layers and proposed their use in Transformer. In a transfer-learning setup, the new variants seem to produce better perplexities for the de-noising objective used in pre-training, as well as better results on many downstream language-understanding tasks. These architectures are simple to implement, and have no apparent computational drawbacks. We offer no explanation as to why these architectures seem to work; we attribute their success, as all else, to divine benevolence.

*GLU Variants Improve Transformer.* Noam Shazeer 2020



# Quick Primer on Compute

It's matmuls all the way down

$$Y = A \cdot B$$

Activations



Weights



$$A \in \mathbb{R}^{M \times D}, B \in \mathbb{R}^{D \times N}$$

We count compute as FLOPs, ie, **F**loating Point **O**perations

Forward FLOPs:  $D$  multiplications and  $D$  additions per entry of  $Y = 2.D.M.N$

Backward FLOPs

**4.D.M.N**

**2x forward pass FLOPs**

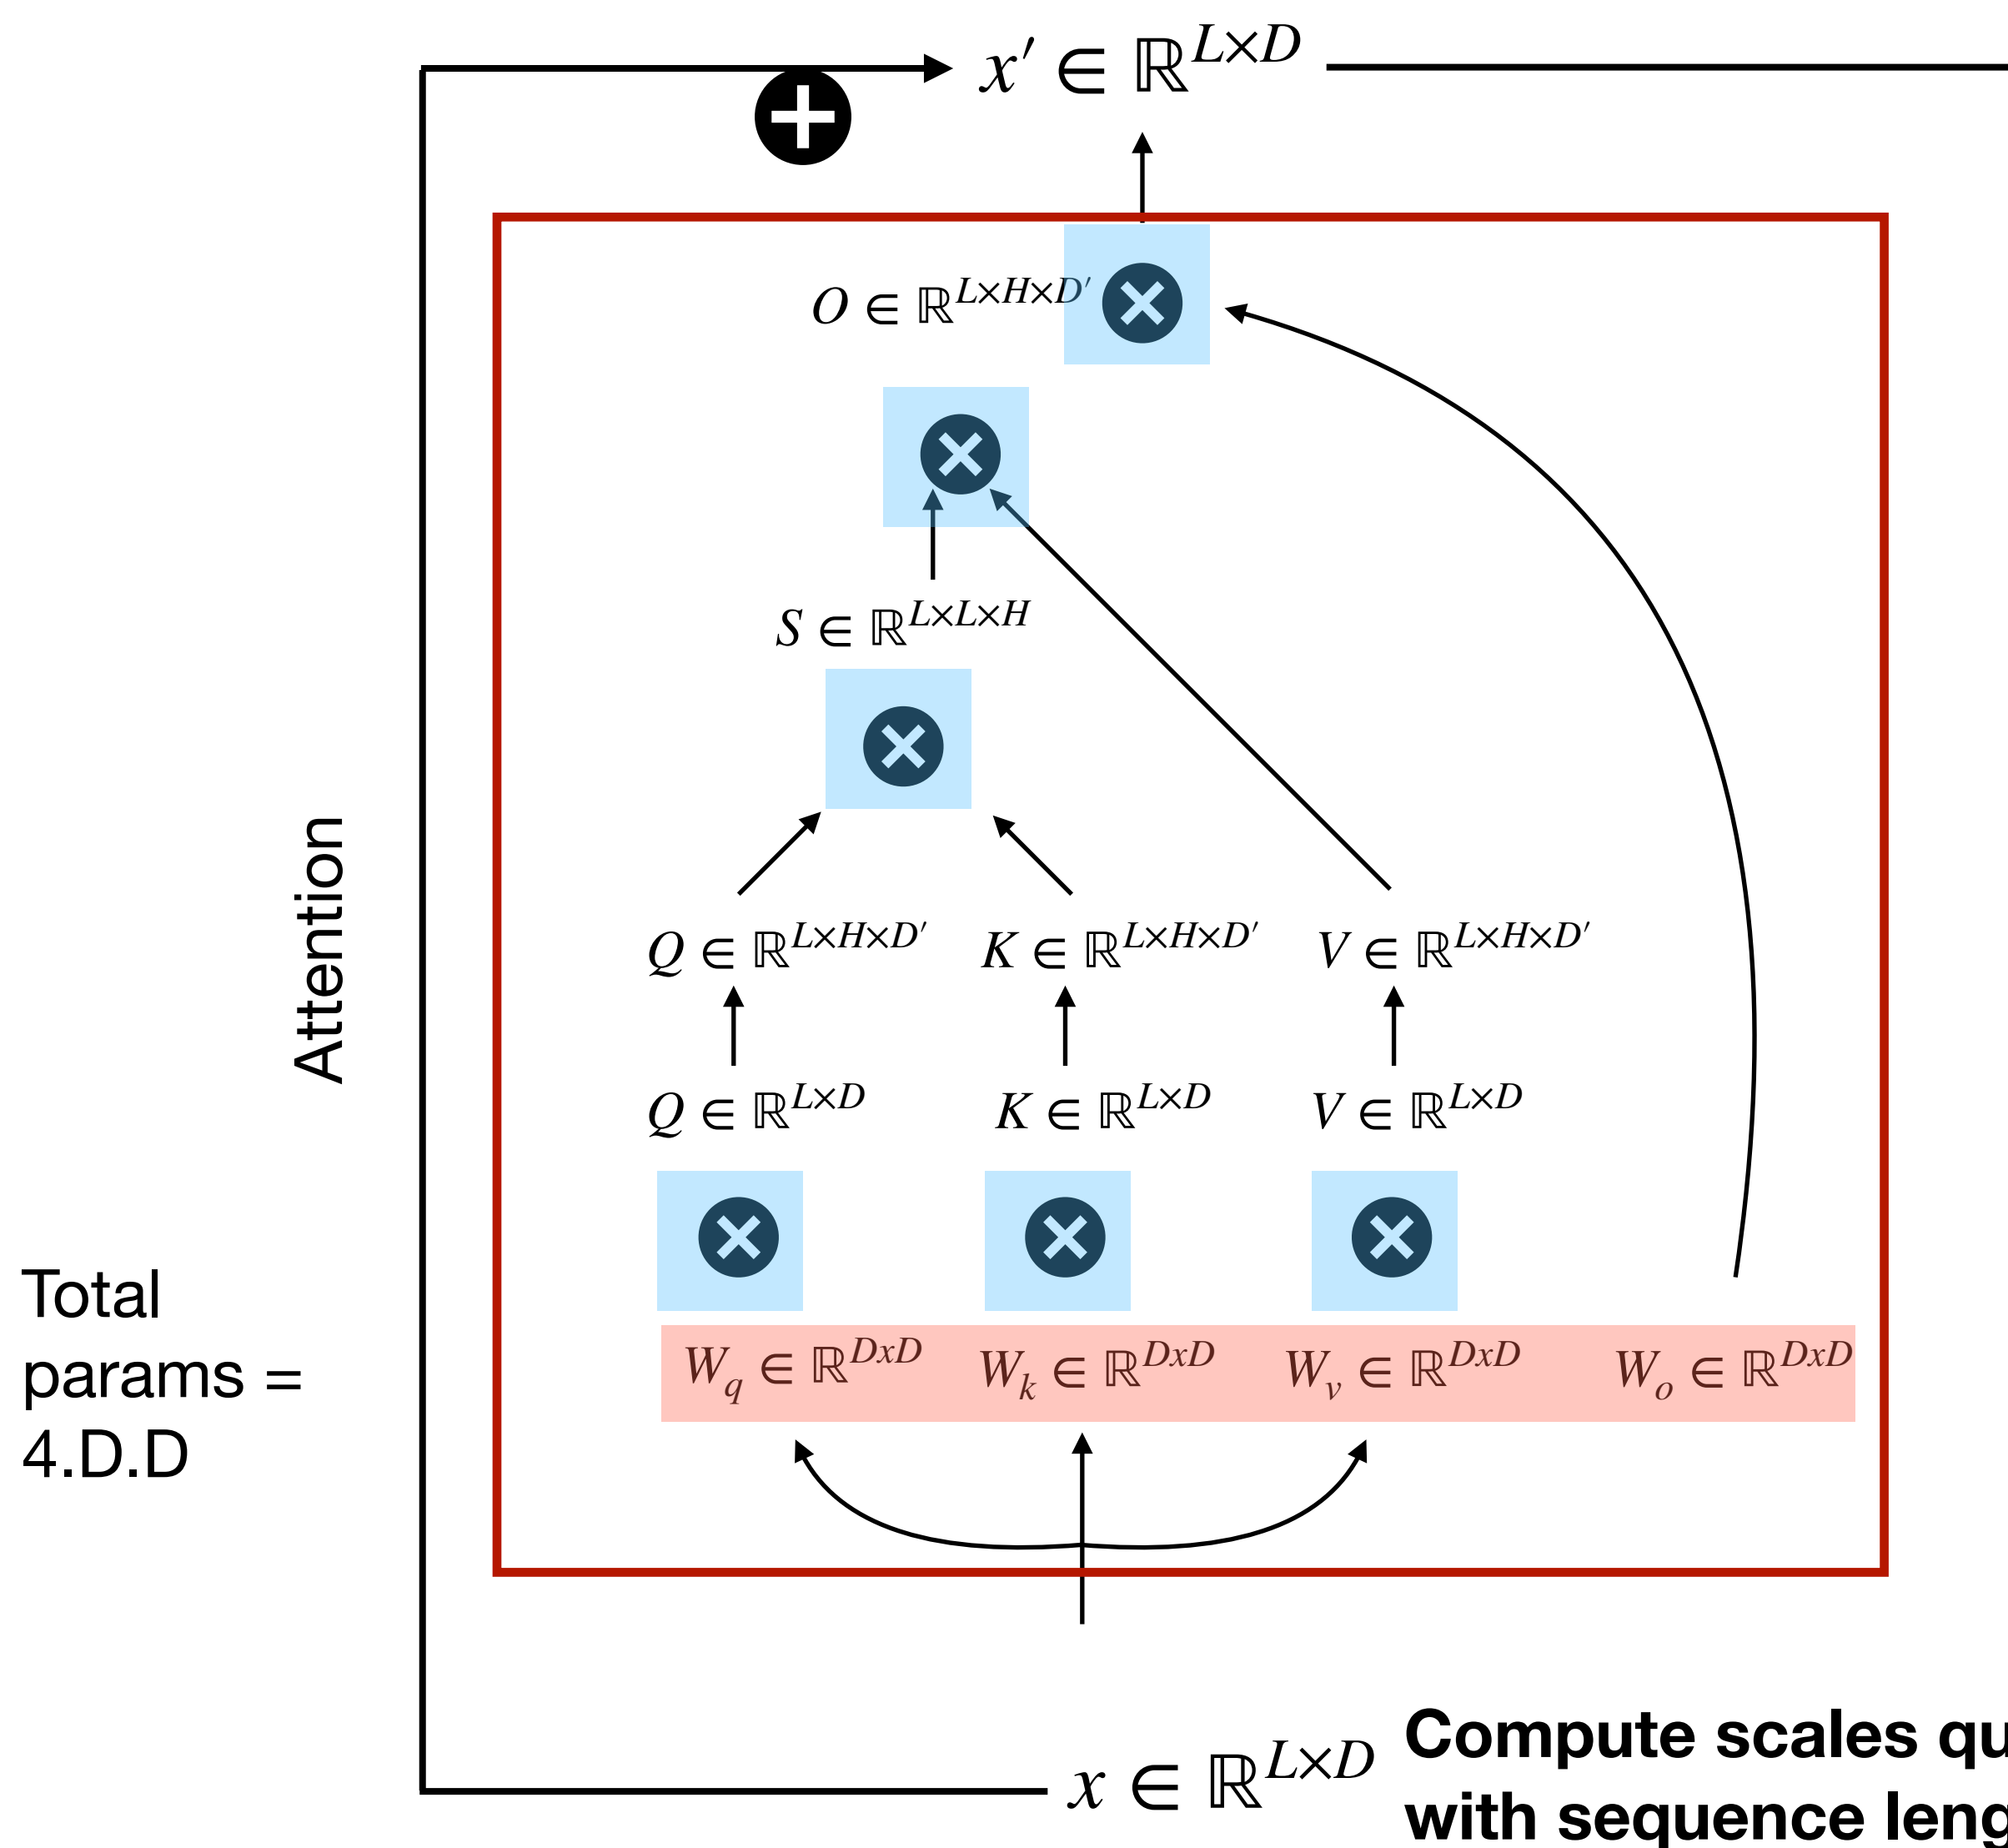
$$\frac{\partial L}{\partial B} = \frac{\partial L}{\partial Y} \cdot \frac{\partial Y}{\partial B} = A^T \cdot \frac{\partial L}{\partial Y}$$

2.D.M.N

$$\frac{\partial L}{\partial A} = \frac{\partial L}{\partial Y} \cdot \frac{\partial Y}{\partial A} = \frac{\partial L}{\partial Y} \cdot B^T$$

2.D.M.N

# Transformers: What Do We Compute?



	FLOPs (inference)	FLOPs (training)	
$x \cdot W_q$	2.L.D.D	6.L.D.D	
$x \cdot W_k$	2.L.D.D	6.L.D.D	
$x \cdot W_v$	2.L.D.D	6.L.D.D	
$Q \cdot K^T$	2.D'.L.L.H = 2.D.L.L	6.D.L.L	
$\text{softmax}$	L (sum every row). L (scale every entry in row). H (do this for every head) = L.L.H	O(L.L.H)	H << D
$S \cdot V$	2.L.L.D'.H = 2.D.L.L	6.D.L.L	
$W_o \cdot O$	2.L.D.D	6.L.D.D	3x more compute in training for same L
	8.L.D.D + 4.D.L.L	24.L.D.D + 12.D.L.L	

# Transformers: What Do We Compute?

	FLOPs (inference)	FLOPs (training)
$x' \cdot W_{in1}$	2.L.D.F	6.L.D.F
$x' \cdot W_{in2}$	2.L.D.F	6.L.D.F
$x' \cdot W_{proj}$	2.L.D.F	6.L.D.F
	6.L.D.F	18.L.D.F

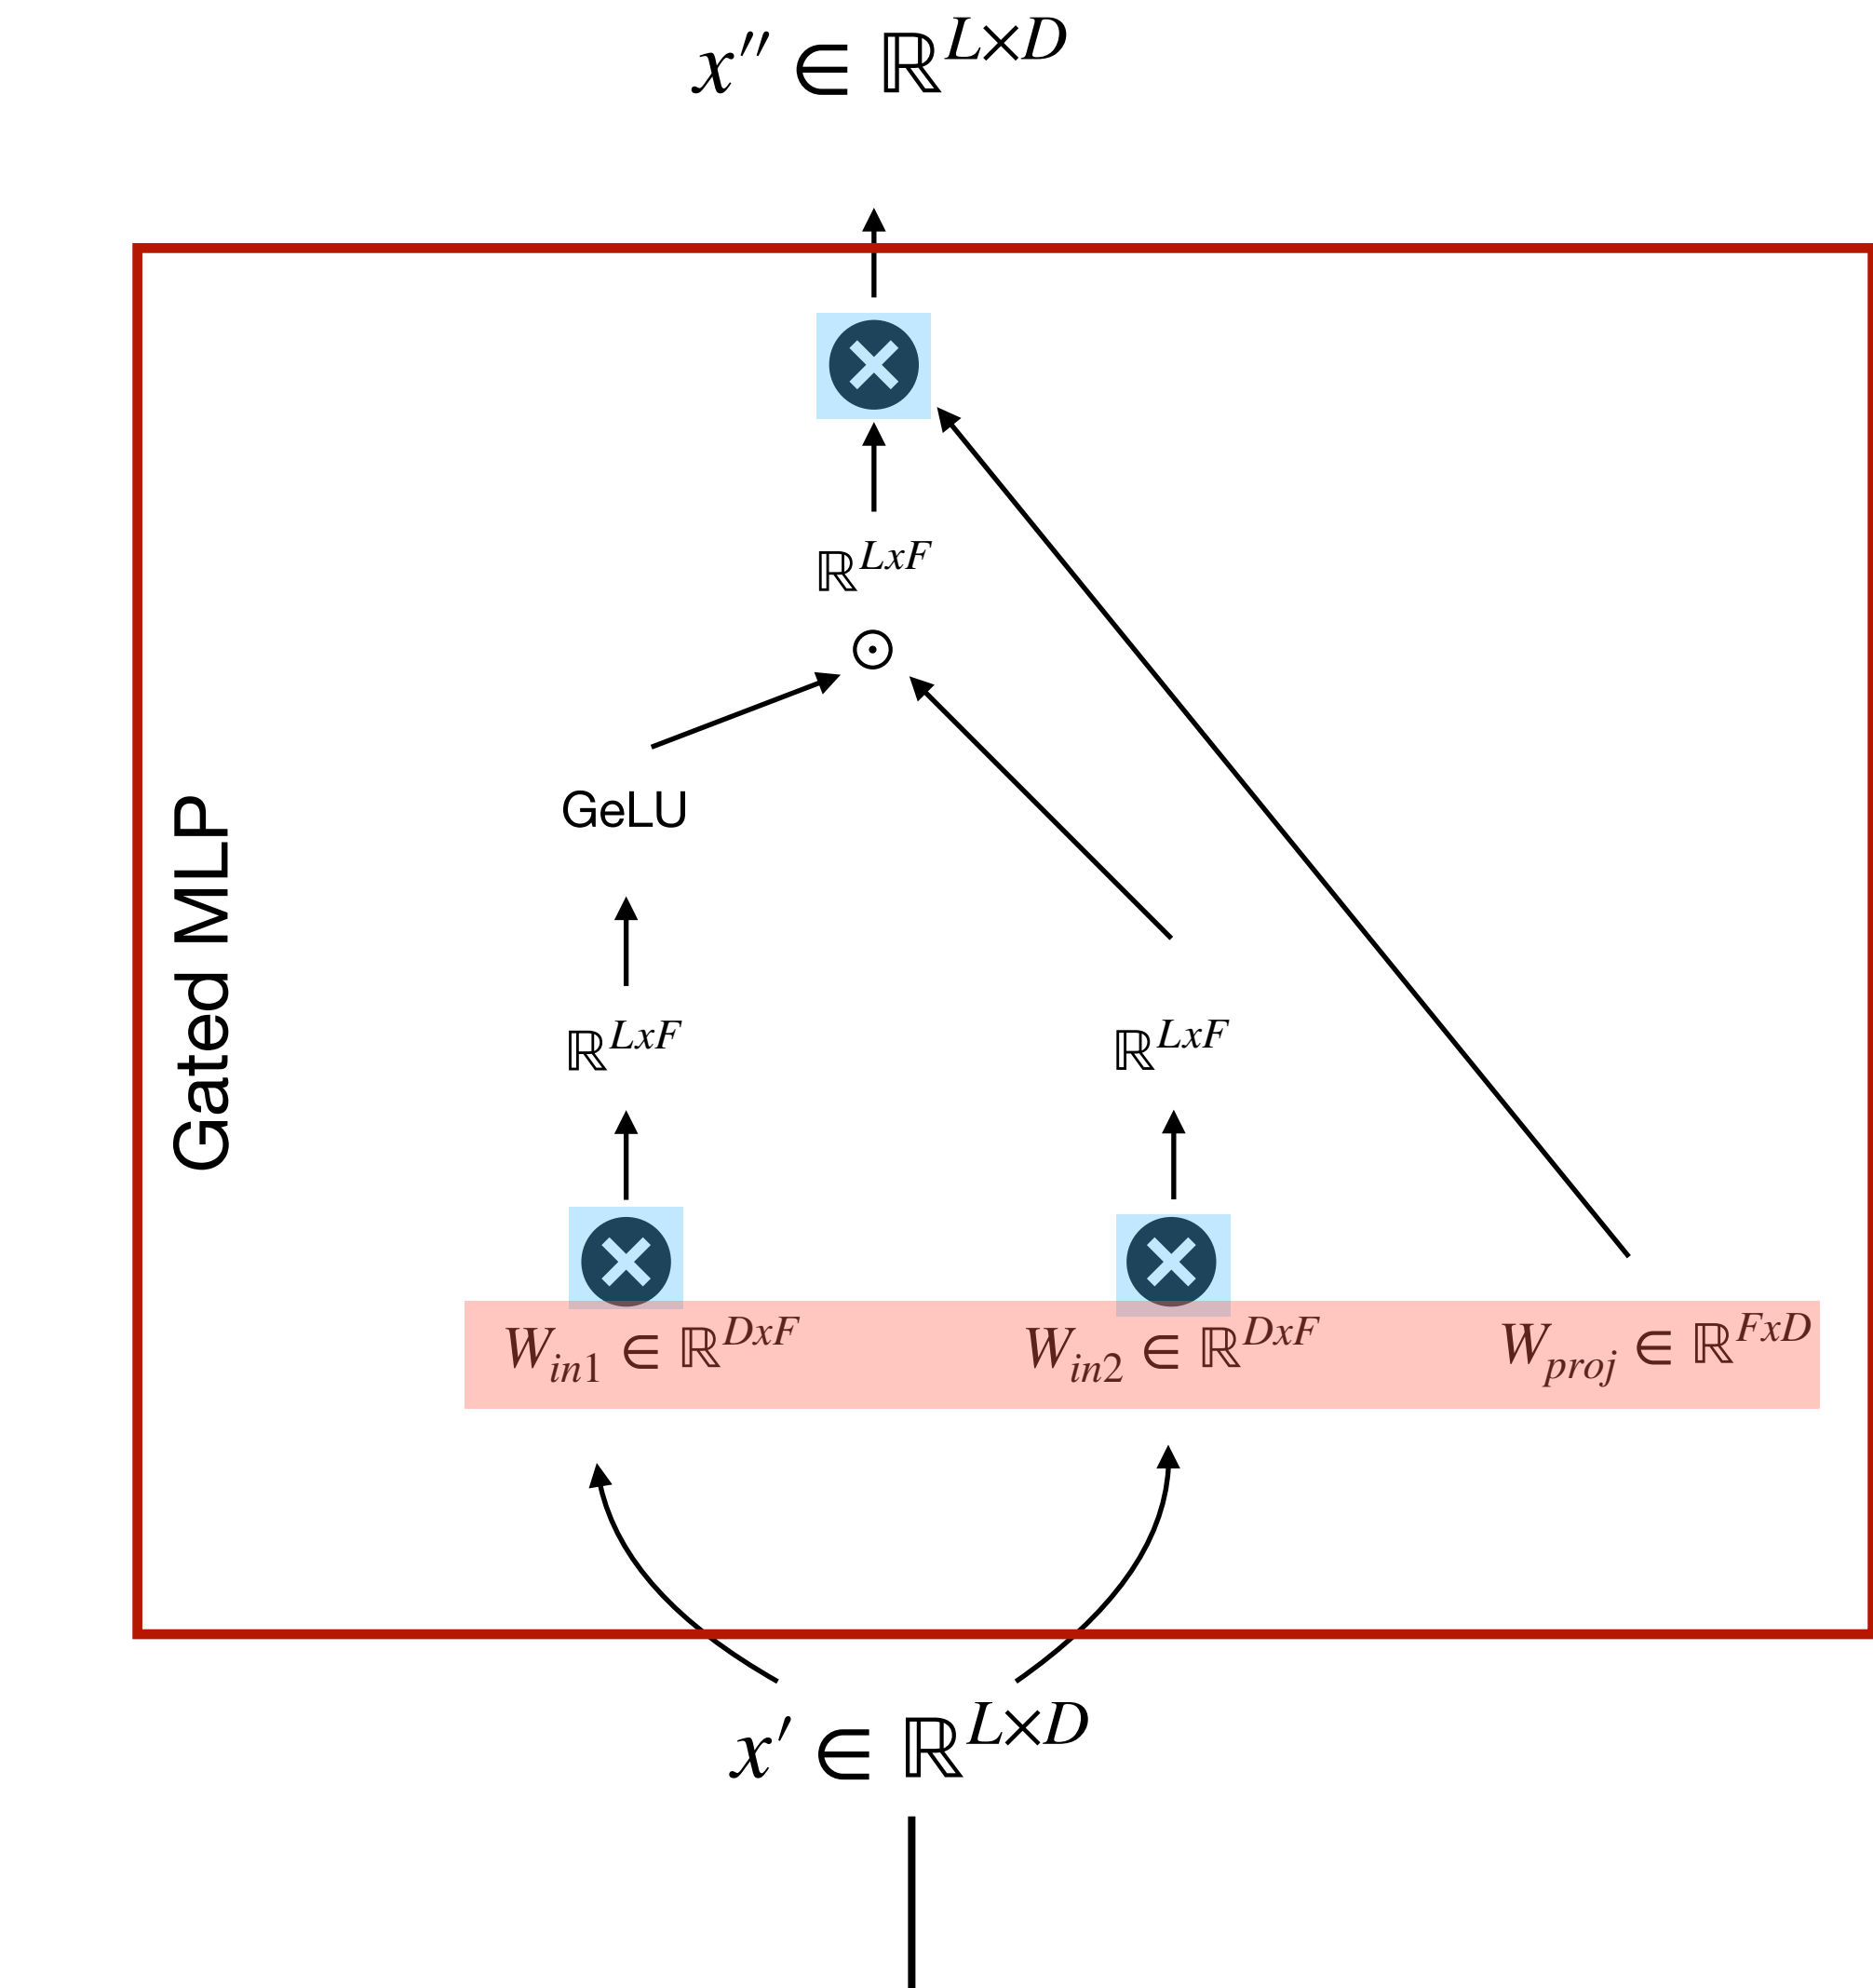
Total Params = 3.D.F

Typically,  $F \approx 3.5D$  (Llama3)

Total Params  $\approx 10.D.D$

vs params in attention = 4.D.D

**Most params in MLP layers!**



# Optional Exercise

Calculate FLOPs for attention variants such as  
GQA (Llama) and MLA (DeepSeek)

# Inference vs Training Compute

For each transformer layer

## Inference

### Attention

$$8.L.D.D + 4.D.L.L$$

### MLP

$$6.L.D.F$$

$$8.L.D.D + 4.D.L.L + 21.L.D.D$$

In the worst case,  $L = O(1)$ ,  
ie one user, one small prompt

$$\text{Total compute} = 29D.D + 4D$$

Extended Reading:  
[Roofline Model](#)

On most GPUs, Inference is memory bound!

## Training

### Attention

$$24.L.D.D + 12.D.L.L$$

### MLP

$$18.L.D.F$$

Plugging  $F = 3.5D$

For most  $L (<4D)$   
MLPs have majority of  
the FLOPs

$$24.L.D.D + 12.D.L.L + 63.L.D.D$$

Typical batch sizes =  $O(1e+6)$  (eg Llama3)

$$\text{Total compute} = 87DD.O(1e+6) + 12D.O(1e+12)$$

Way more compute in training!

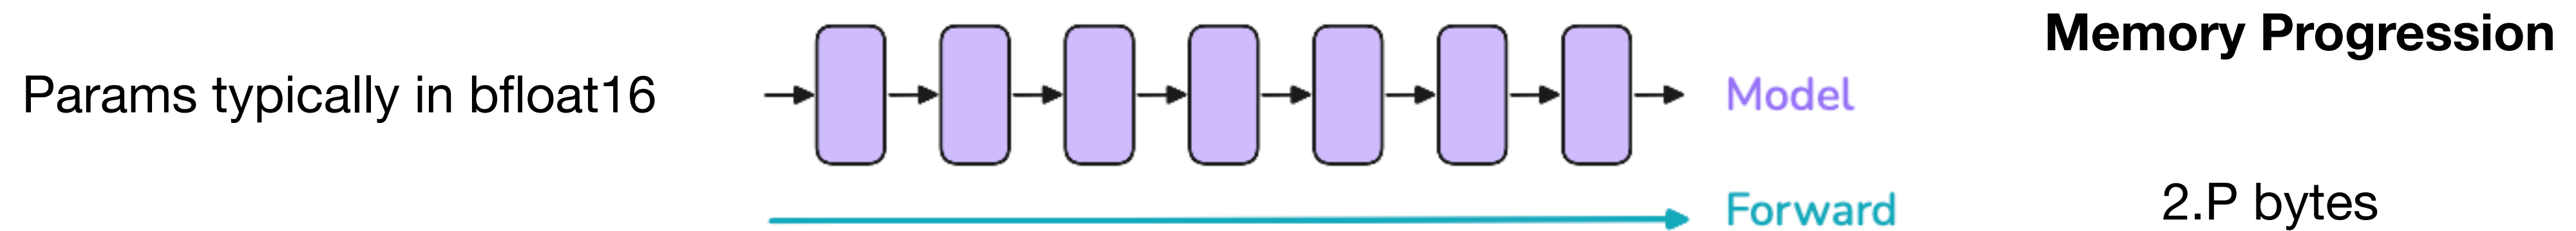
# Profiling Inference and Training

Hands-on exercise: Use PyTorch profiler to analyze inference vs training workloads



# LLM Training

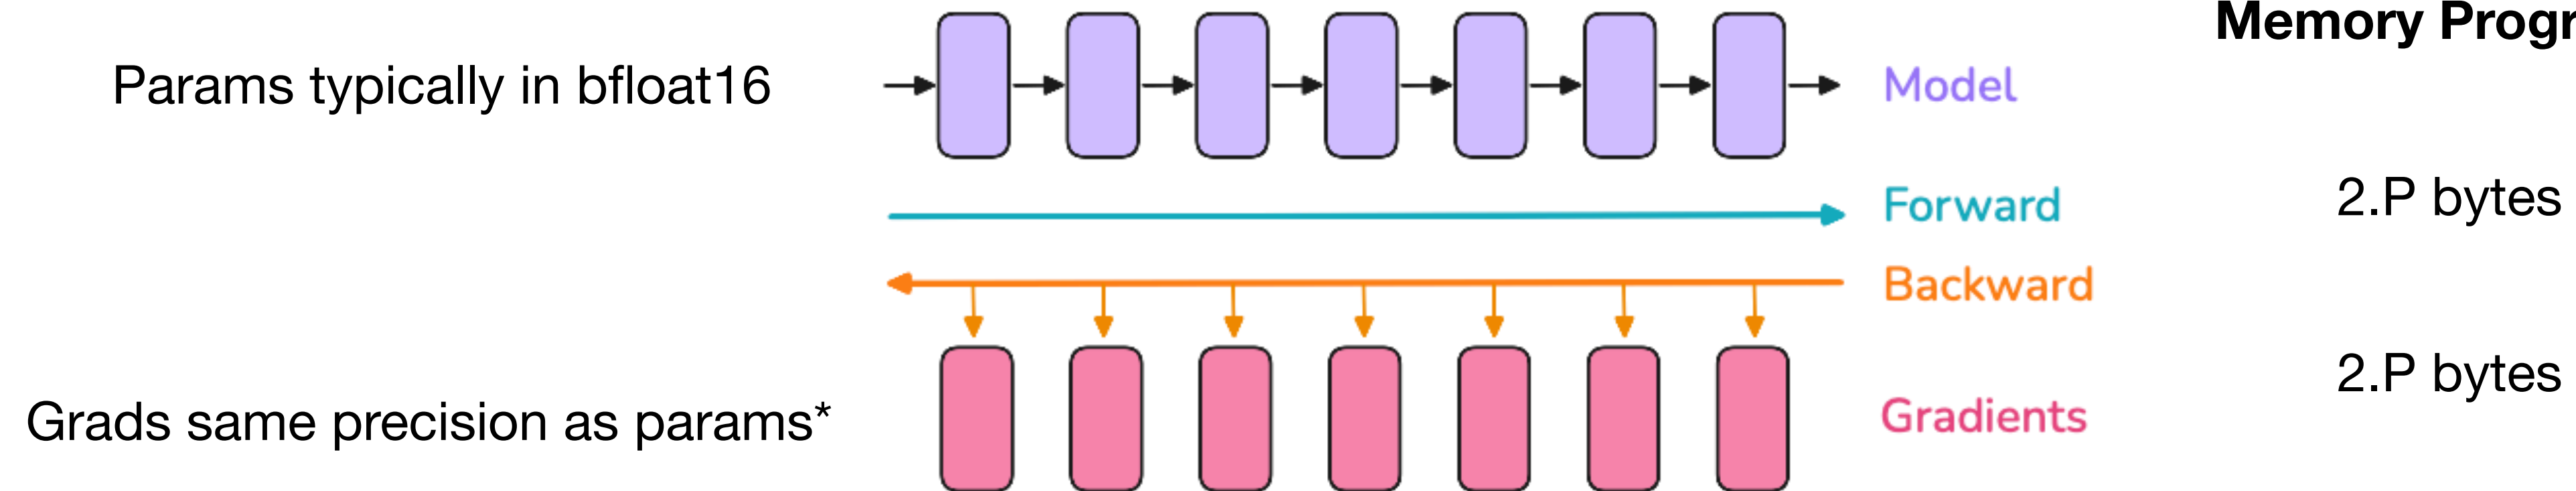
*Key difference from inference: We hold activations in memory for the backward pass*



# LLM Training

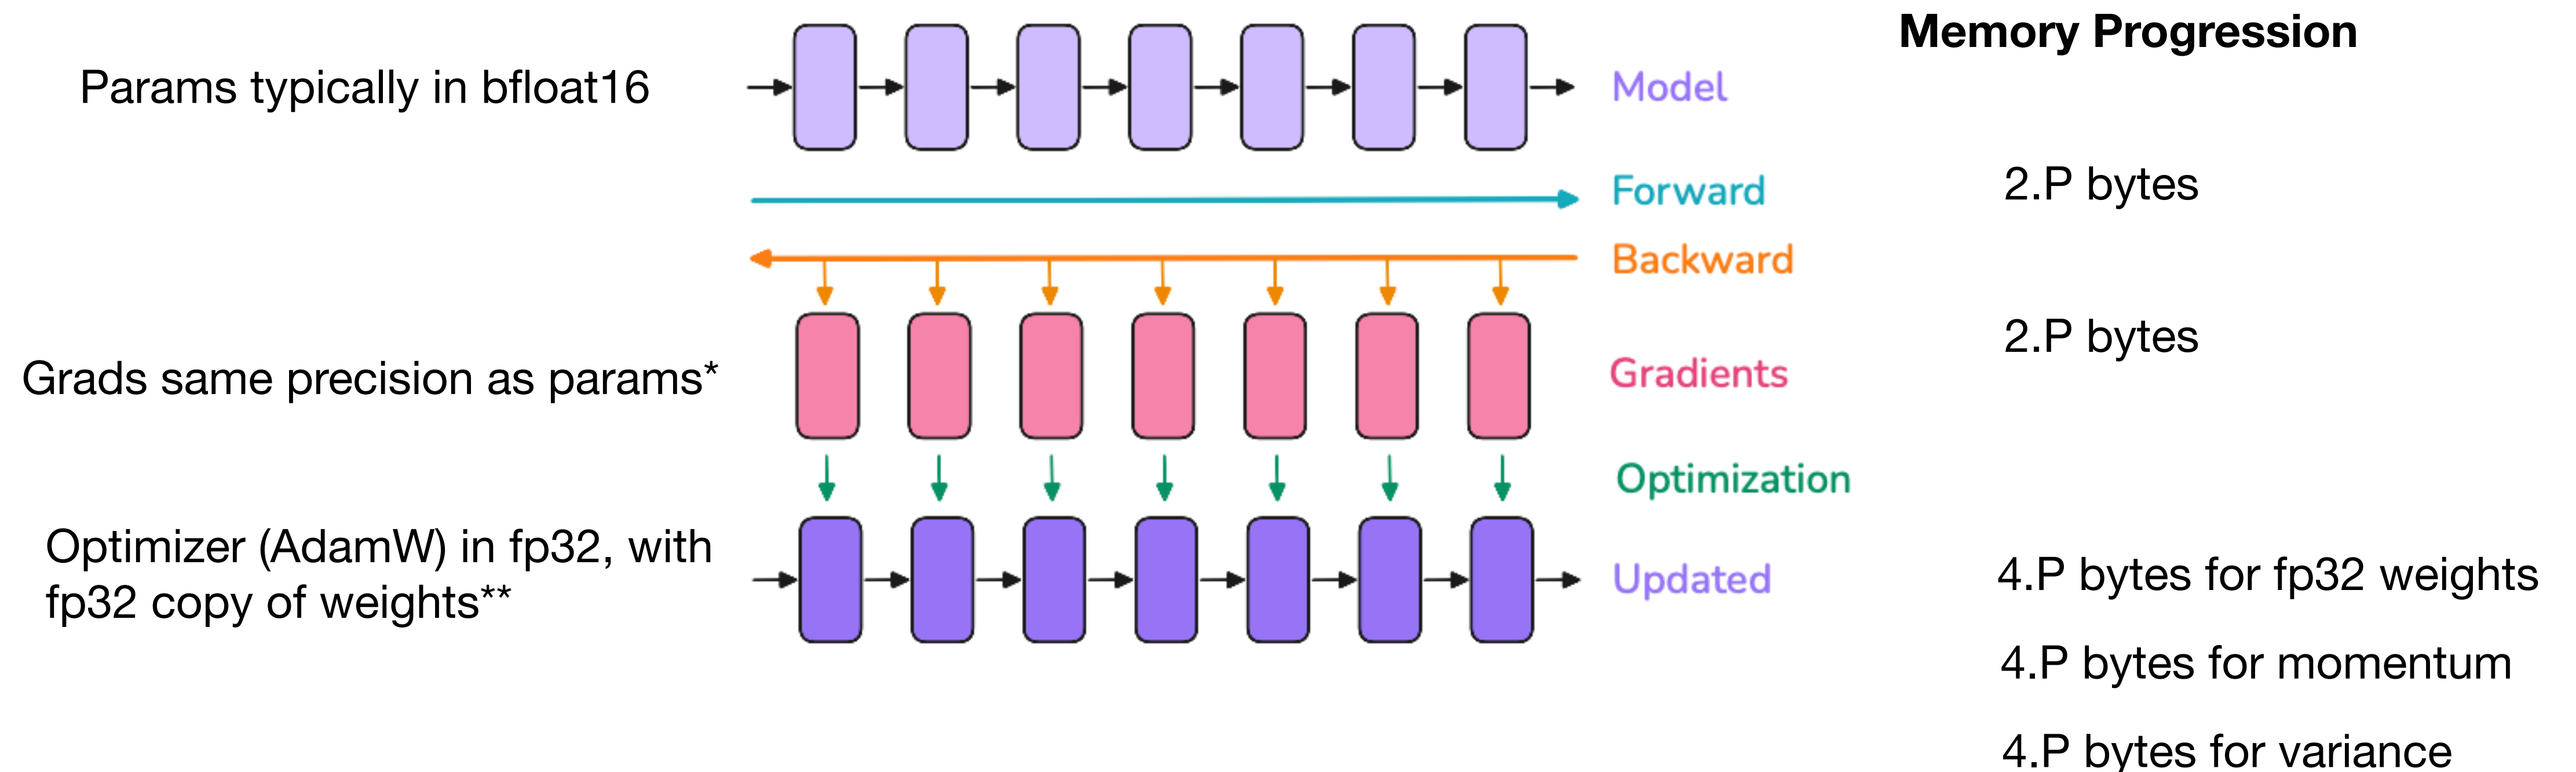
*Key difference from inference: We hold activations in memory for the backward pass*

## Memory Progression



# LLM Training

*Key difference from inference: We hold activations in memory for the backward pass*

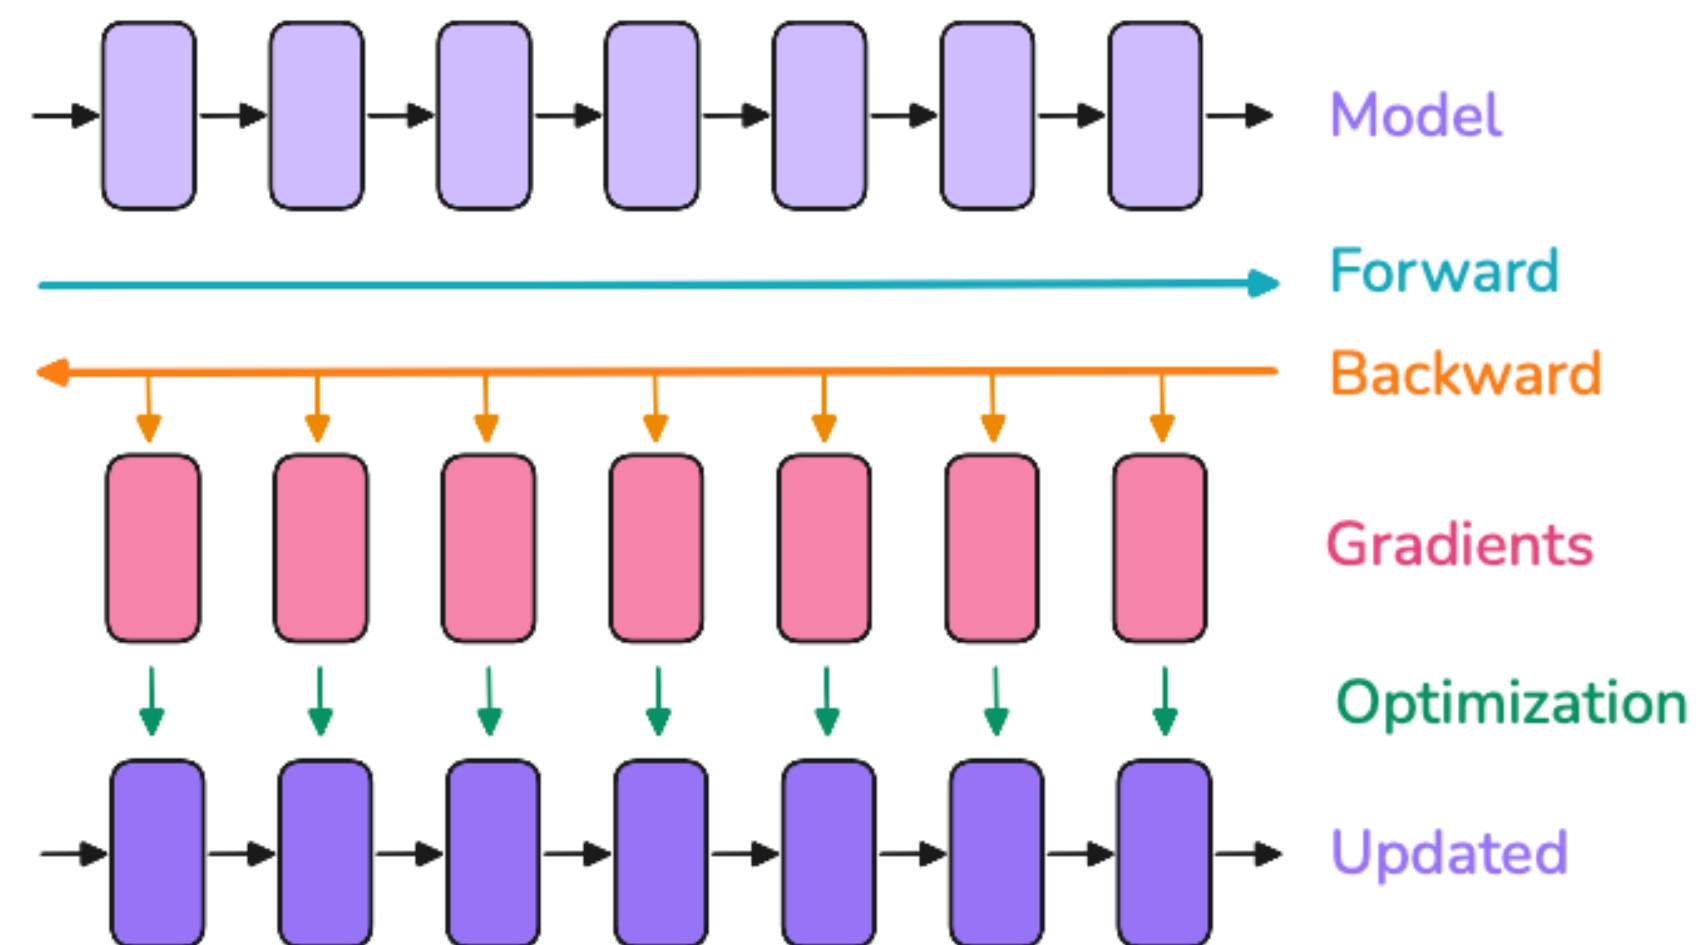


*Ultrascale Playbook. HuggingFace 2025*

\* Later we will see that sometimes, for stability it makes sense to store gradients in fp32

\*\* As a rule of thumb, operations that “accumulate” should be done in fp32 to avoid underflow in bf16

# LLM Training: Memory



	Mixed precision	Mixed precision w fp32 grads	Full precision
Params	2.P	2.P	4.P
Grads	2.P	4.P	4.P
Optimizer	12.P	12.P	8.P
	16.P	18.P	16.P
Llama 8B	128GB	144GB	128GB

Q1: Why use mixed precision when it requires same memory as full?

Q2: Are we missing something?

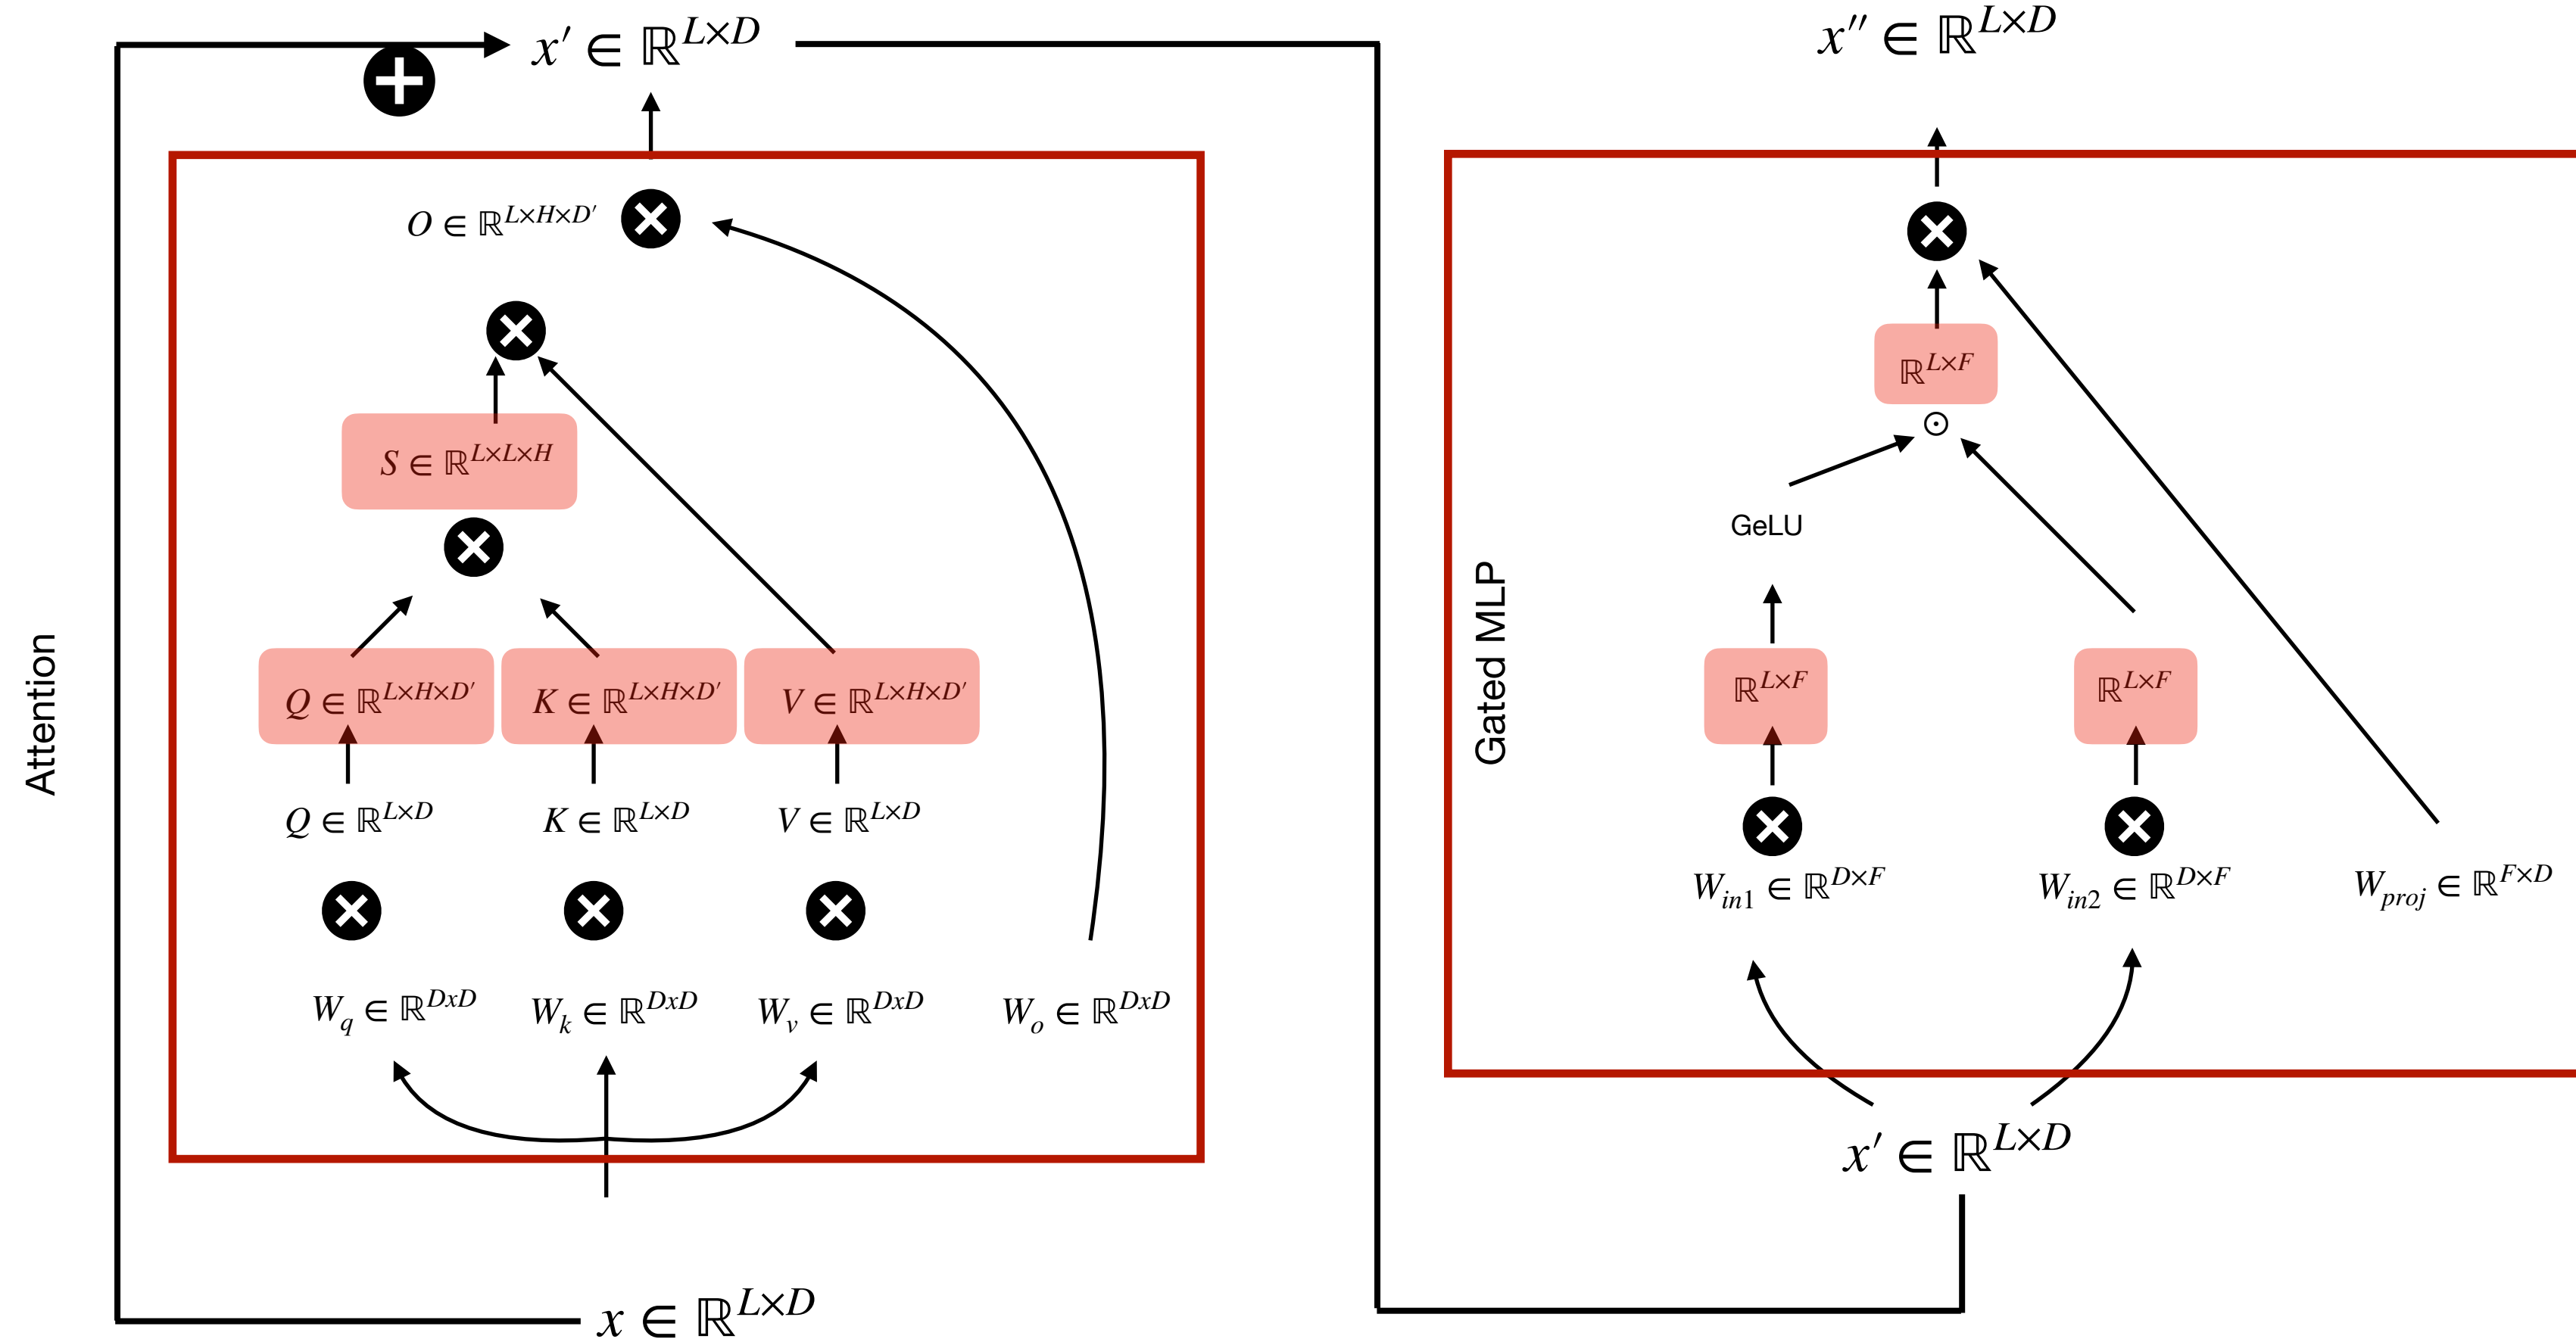
**Activations!**

*Recall: We hold activations in memory for the backward pass*

Technical Specifications		
	H100 SXM	H100 NVL
FP64	34 teraFLOPS	30 teraFLOPS
FP64 Tensor Core	67 teraFLOPS	60 teraFLOPS
FP32	67 teraFLOPS	60 teraFLOPS
TF32 Tensor Core*	989 teraFLOPS	835 teraFLOPS
BFLOAT16 Tensor Core*	1,979 teraFLOPS	1,671 teraFLOPS

**Much  
faster!**

# How Much Do Activations Need?



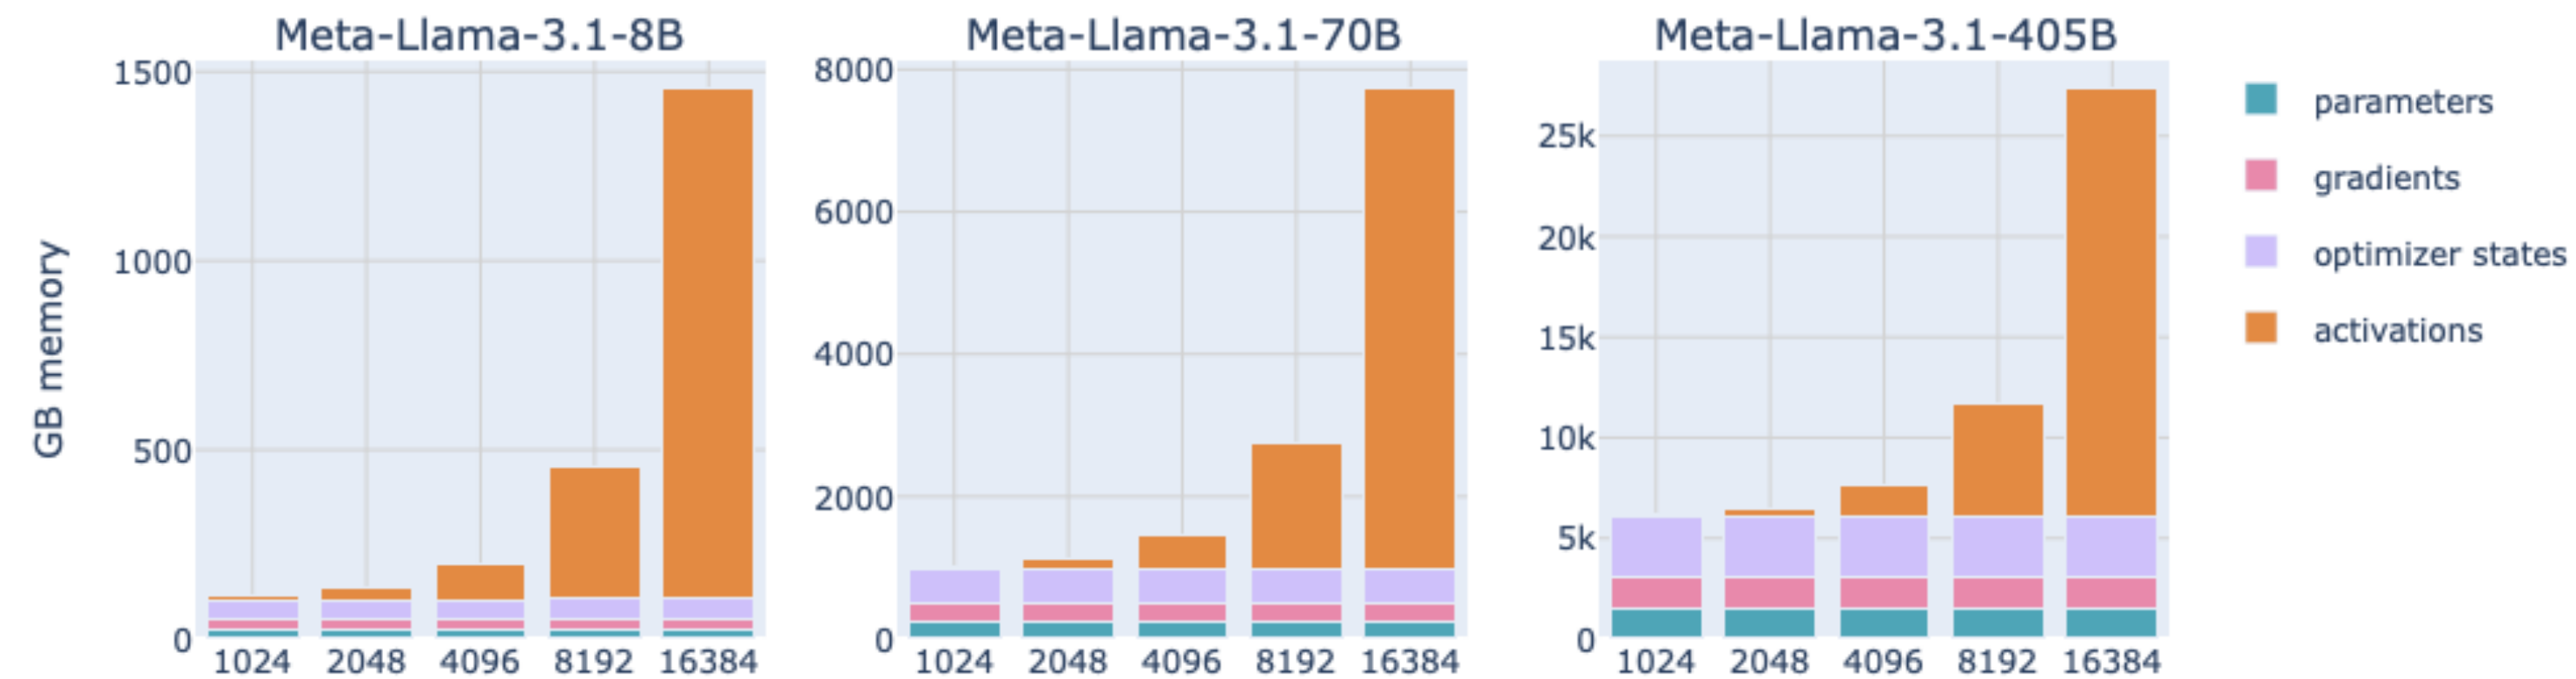
Major contributor:  
materializing the  
attention matrix

For 1 sequence of length  
 $L$ , act mem =  $O(L^2)$

For `bs` sequences of length  
 $L$ , act mem =  $O(\text{bs} \cdot L^2)$

# How Does Memory Scale?

Keeping bs = 1, for different Llama model sizes



For smaller sequences, this is negligible, but it very quickly dwarfs the memory footprint of params, grads, optimizer combined!

# Recap

- Calculating FLOPs in transformer layers for inference vs training
- Understanding compute and memory in LLM training workloads



# Fitting Things on One GPU

Mixed precision

Params 2.P

Grads 2.P

Optimizer 12.P

16.P

Llama 8B 128GB



You cannot train a model where params + grads + optimizer don't fit in memory

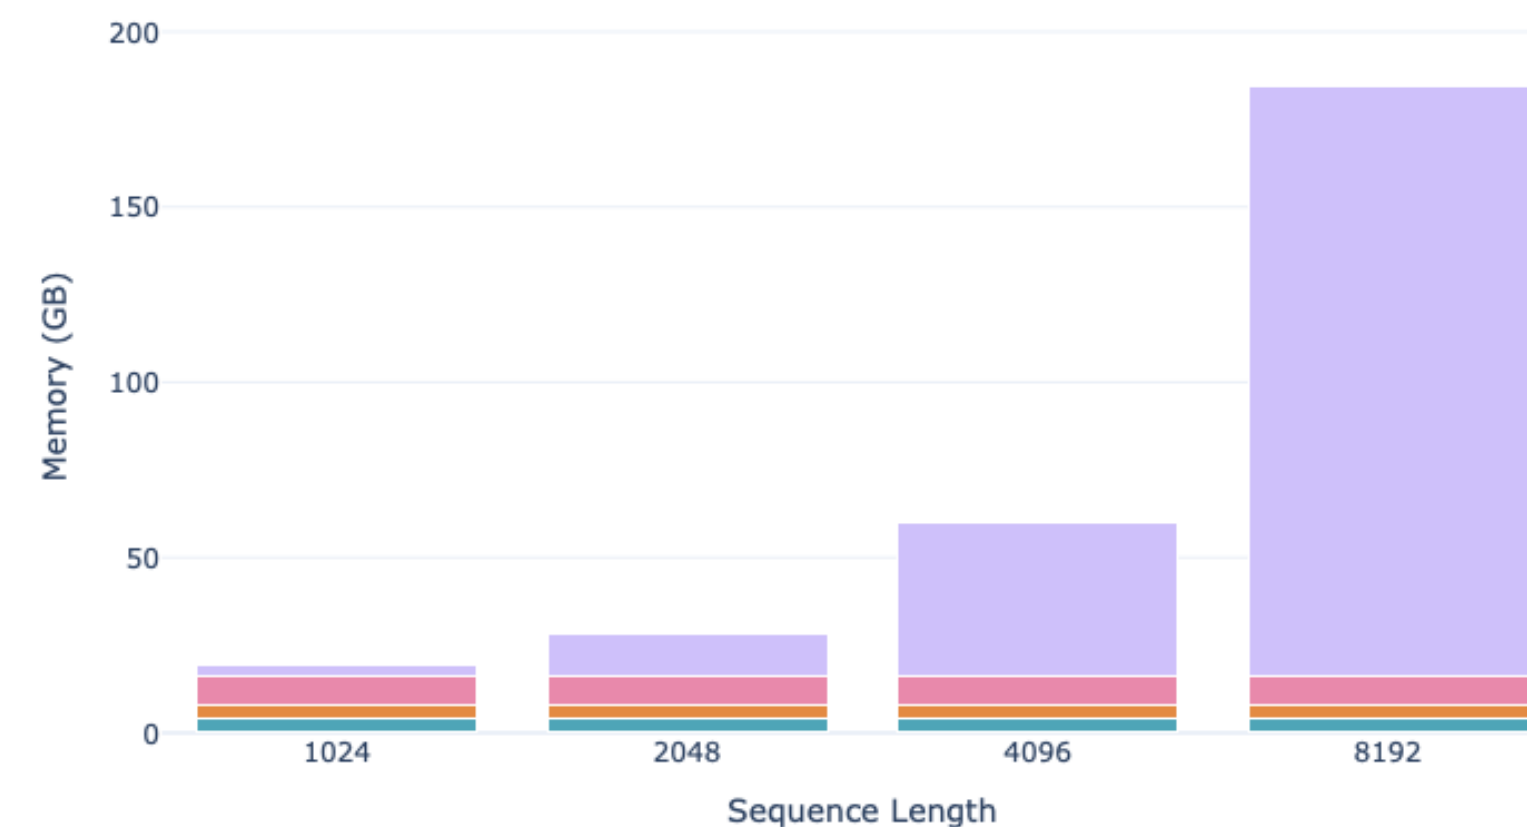
=> Reduce parameter count

Llama 1B

16GB



However...



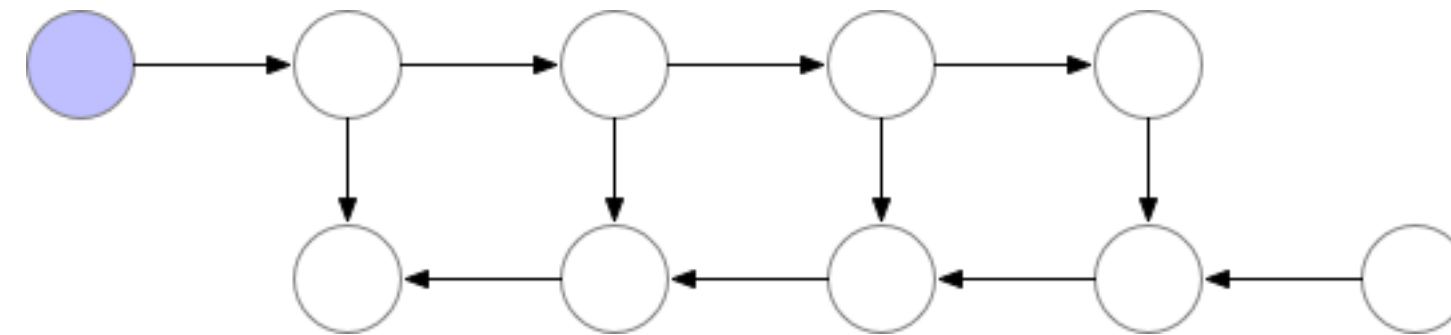
# Activation Checkpointing

Key concept: Trade off memory for compute

Other names: gradient checkpointing, activation recomputation, rematerialization

*Drop activations from memory and re-compute when needed*

Default

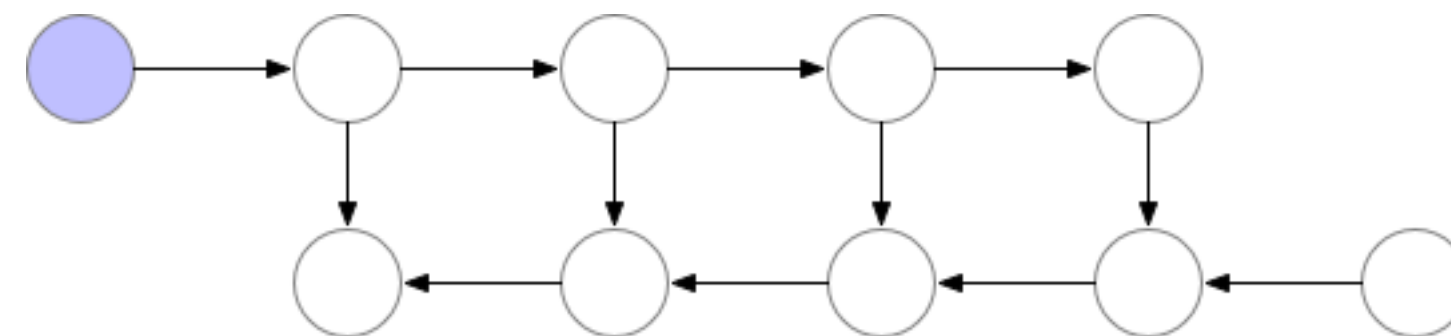


Memory:  $O(\text{num\_layers})$

Compute:  $O(\text{num\_layers})$

Forward per layer: 1

Keep every  $\sqrt{\text{num\_layer}}$   
in memory



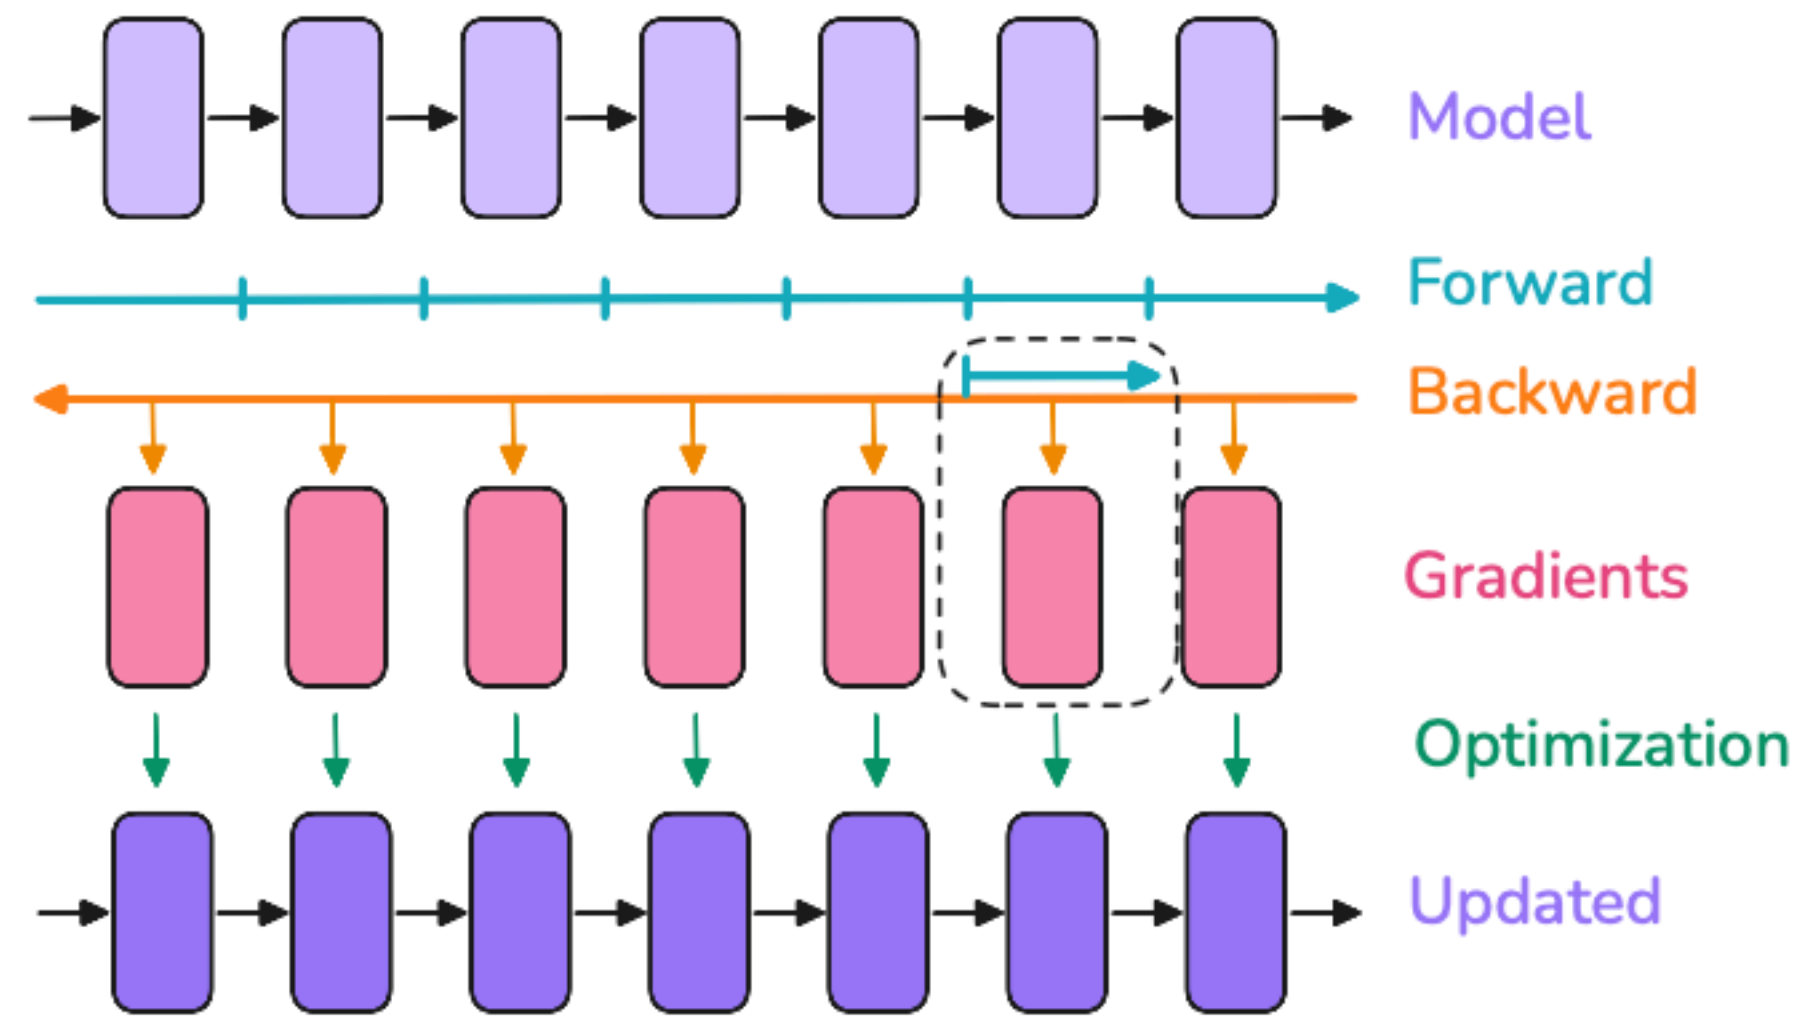
Memory:  $O(\sqrt{\text{num\_layers}})$

Compute:  $O(\text{num\_layers})$

Forward per layer: 1 to 2

# Activation Checkpointing

Easy to implement in PyTorch



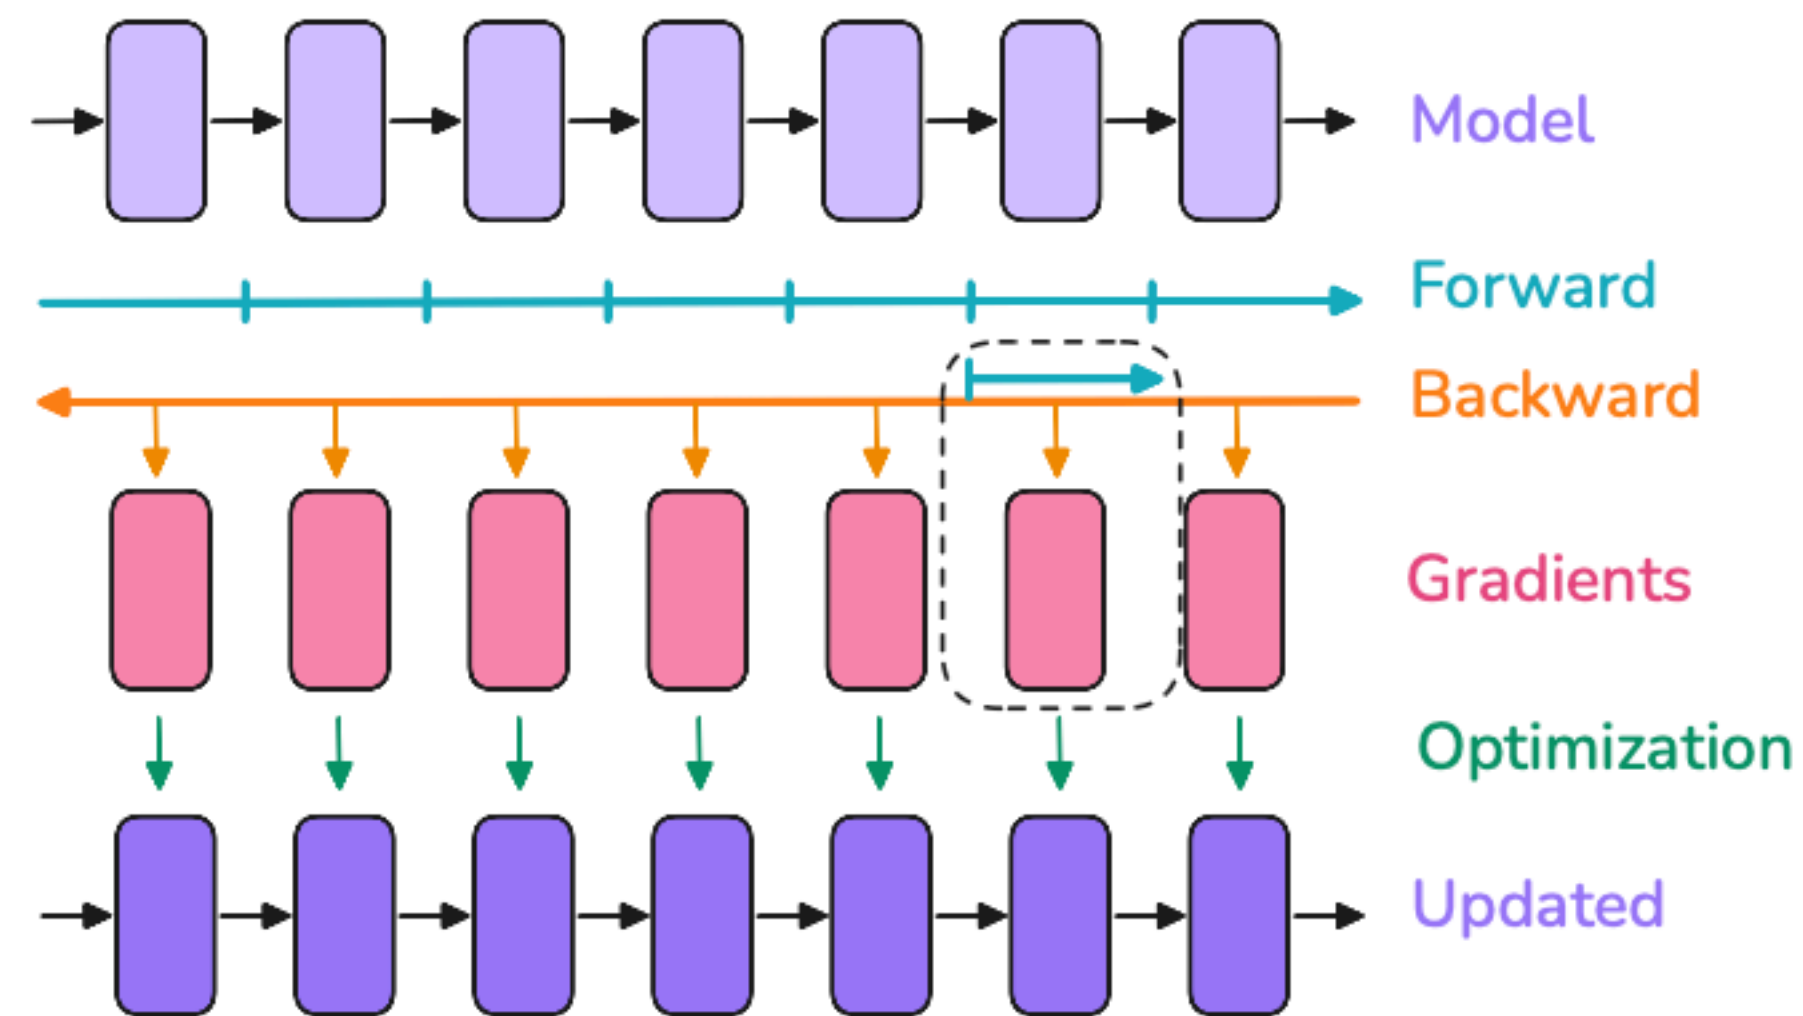
```
class Model(nn.Module):  
    def __init__(self, input_size=512, hidden_size=1024, num_layers=8): ...  
  
    def forward(self, x):  
        for layer in self.layers:  
            x = layer(x)  
        return x
```

```
from torch.utils.checkpoint import checkpoint  
# Model with Activation Checkpointing on Every Layer  
class CheckpointedModel(nn.Module):  
    def __init__(self, input_size=512, hidden_size=1024, num_layers=8, use_checkpoint=True): ...  
  
    def forward(self, x):  
        for layer in self.layers:  
            # Use checkpointing during training to save memory  
            # The checkpoint function will:  
            # 1. Run the forward pass normally  
            # 2. Discard intermediate activations  
            # 3. Recompute them during backward pass when needed  
            x = checkpoint(layer, x, use_reentrant=False)  
  
        return x
```

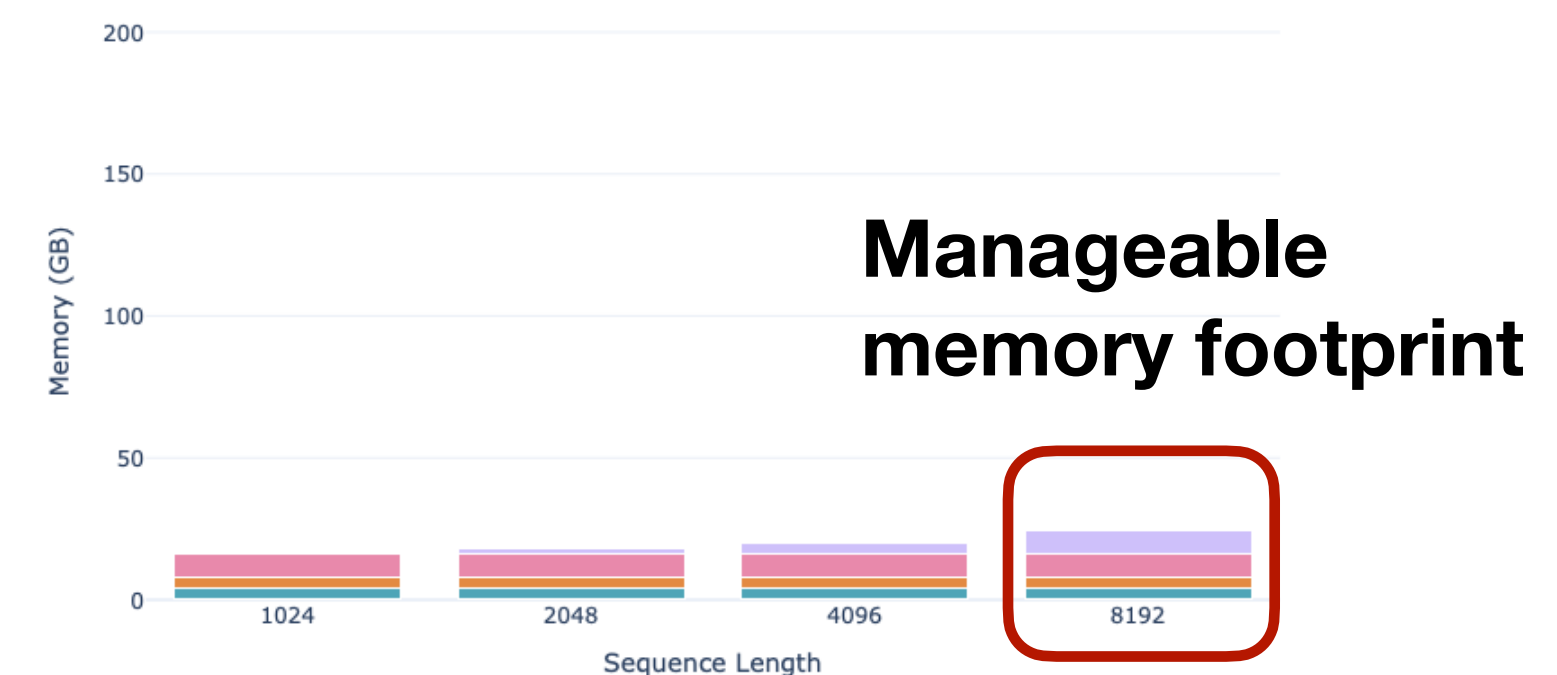
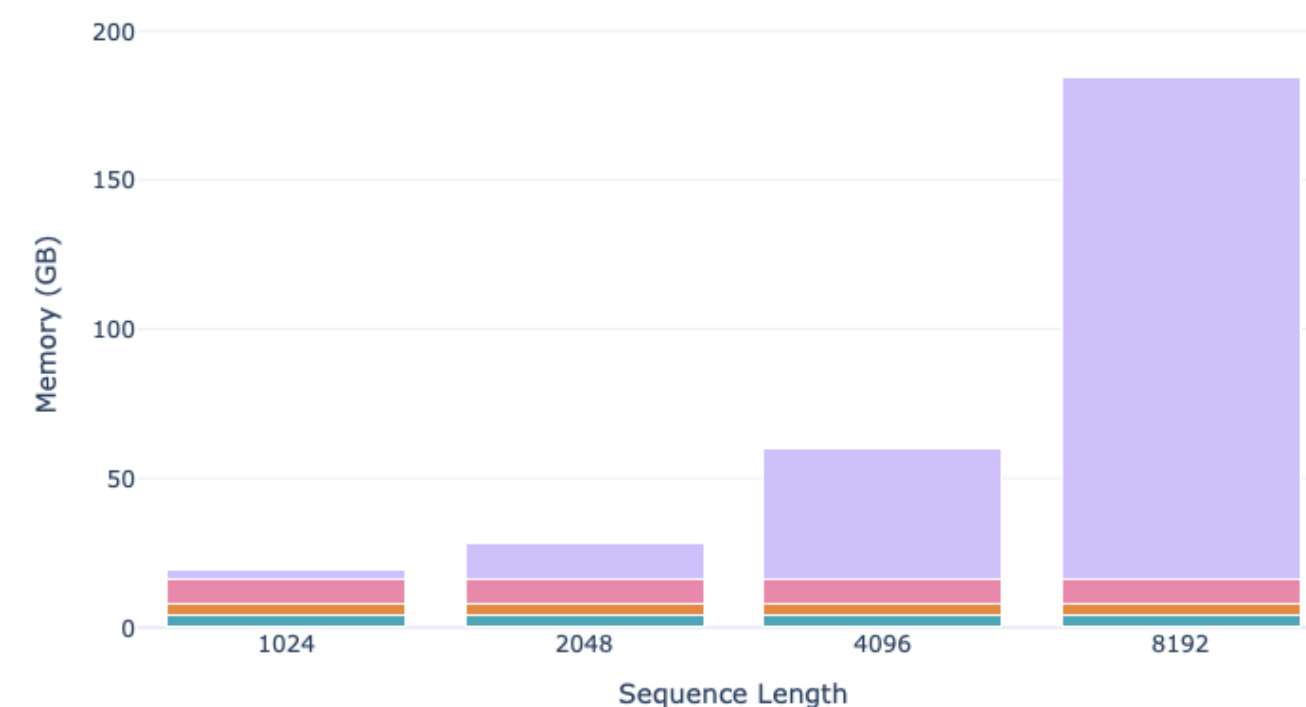
<https://docs.pytorch.org/docs/stable/checkpoint.html>

# Activation Checkpointing

*Drop activations from memory and re-compute when needed*



1. Flash Attention comes with in-built recomputation — takes care of exponential explosion from materializing attention matrix
2. Selective re-computation (i.e. checkpointing after residual of a transformer block) is generally a good rule of thumb
3. FLOPs  $\uparrow$  , Memory  $\downarrow$



# Are we done?

*Recall: we kept  $bs=1$ , and scaled sequence length*

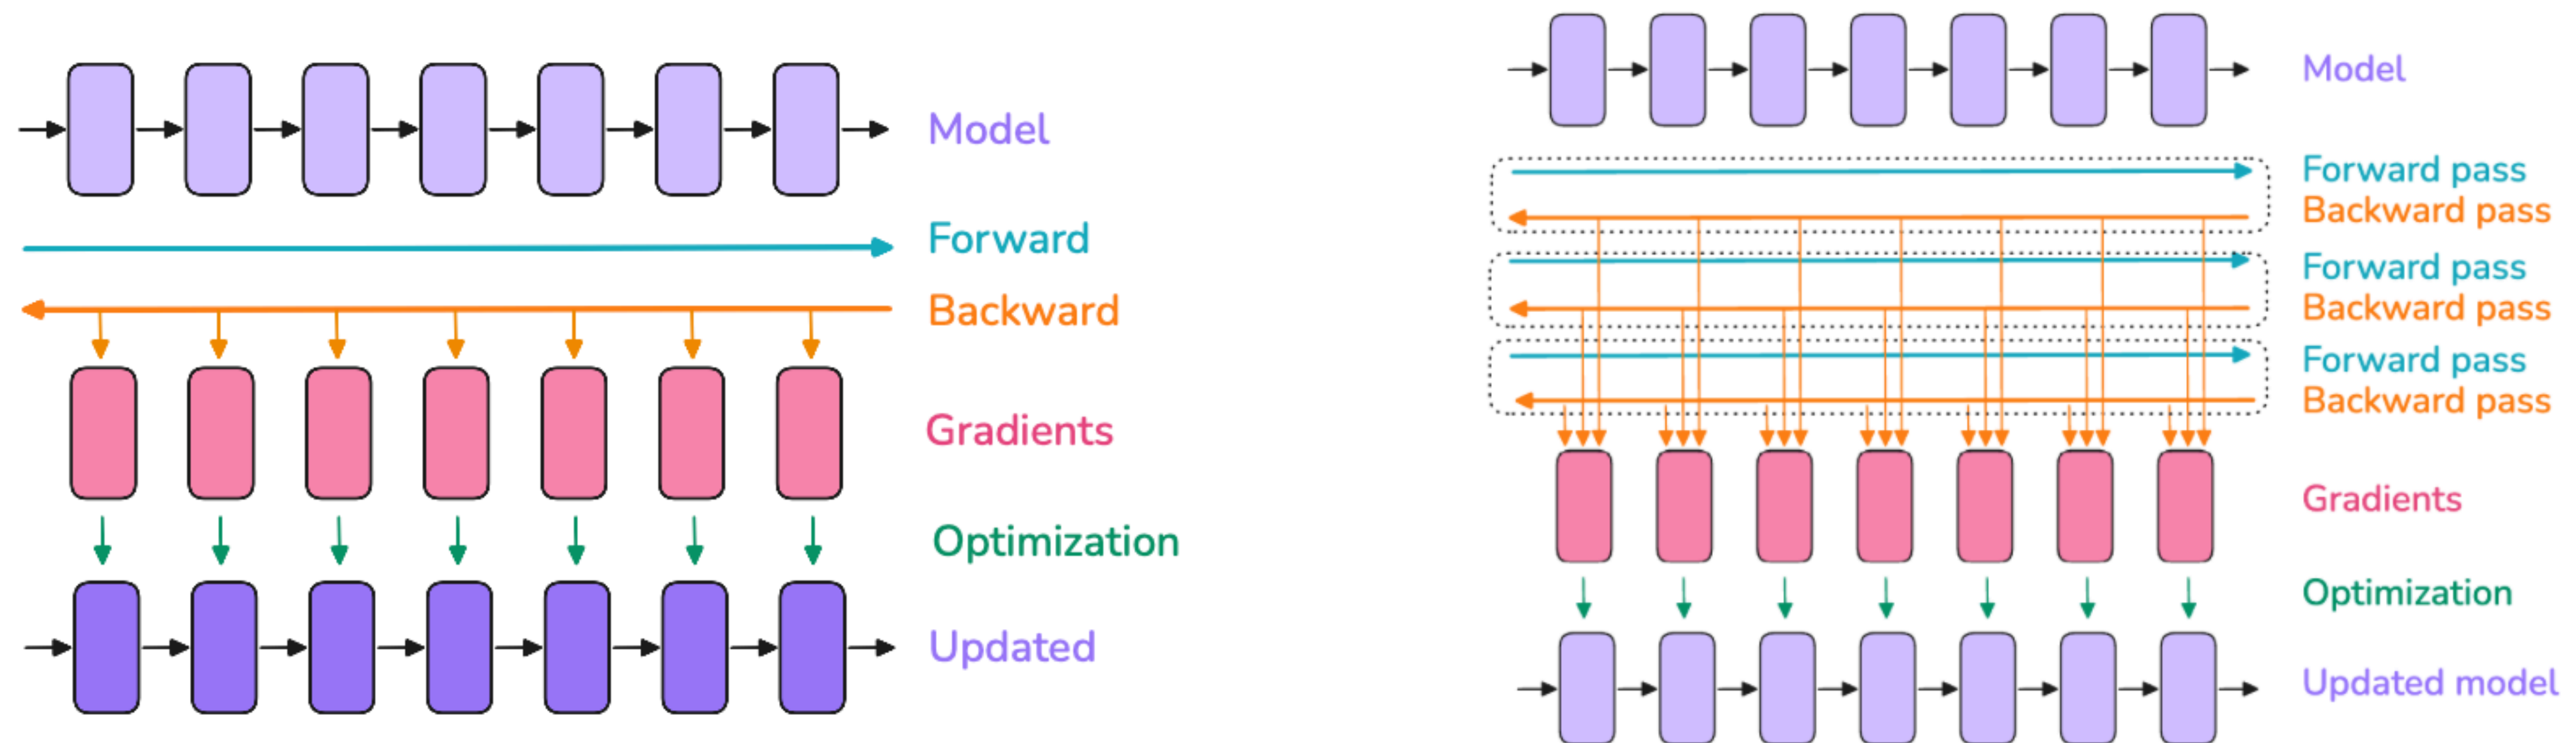
*Also recall: For `bs` sequences of length  $L$ , act mem =  $O(bs.L^2)$*

**How do we scale bs without exploding activation memory?**

# Gradient Accumulation

*Problem: we wish to emulate a larger bs than memory permits*

*Key observation: we don't have to immediately do an optimizer step after computing gradients*

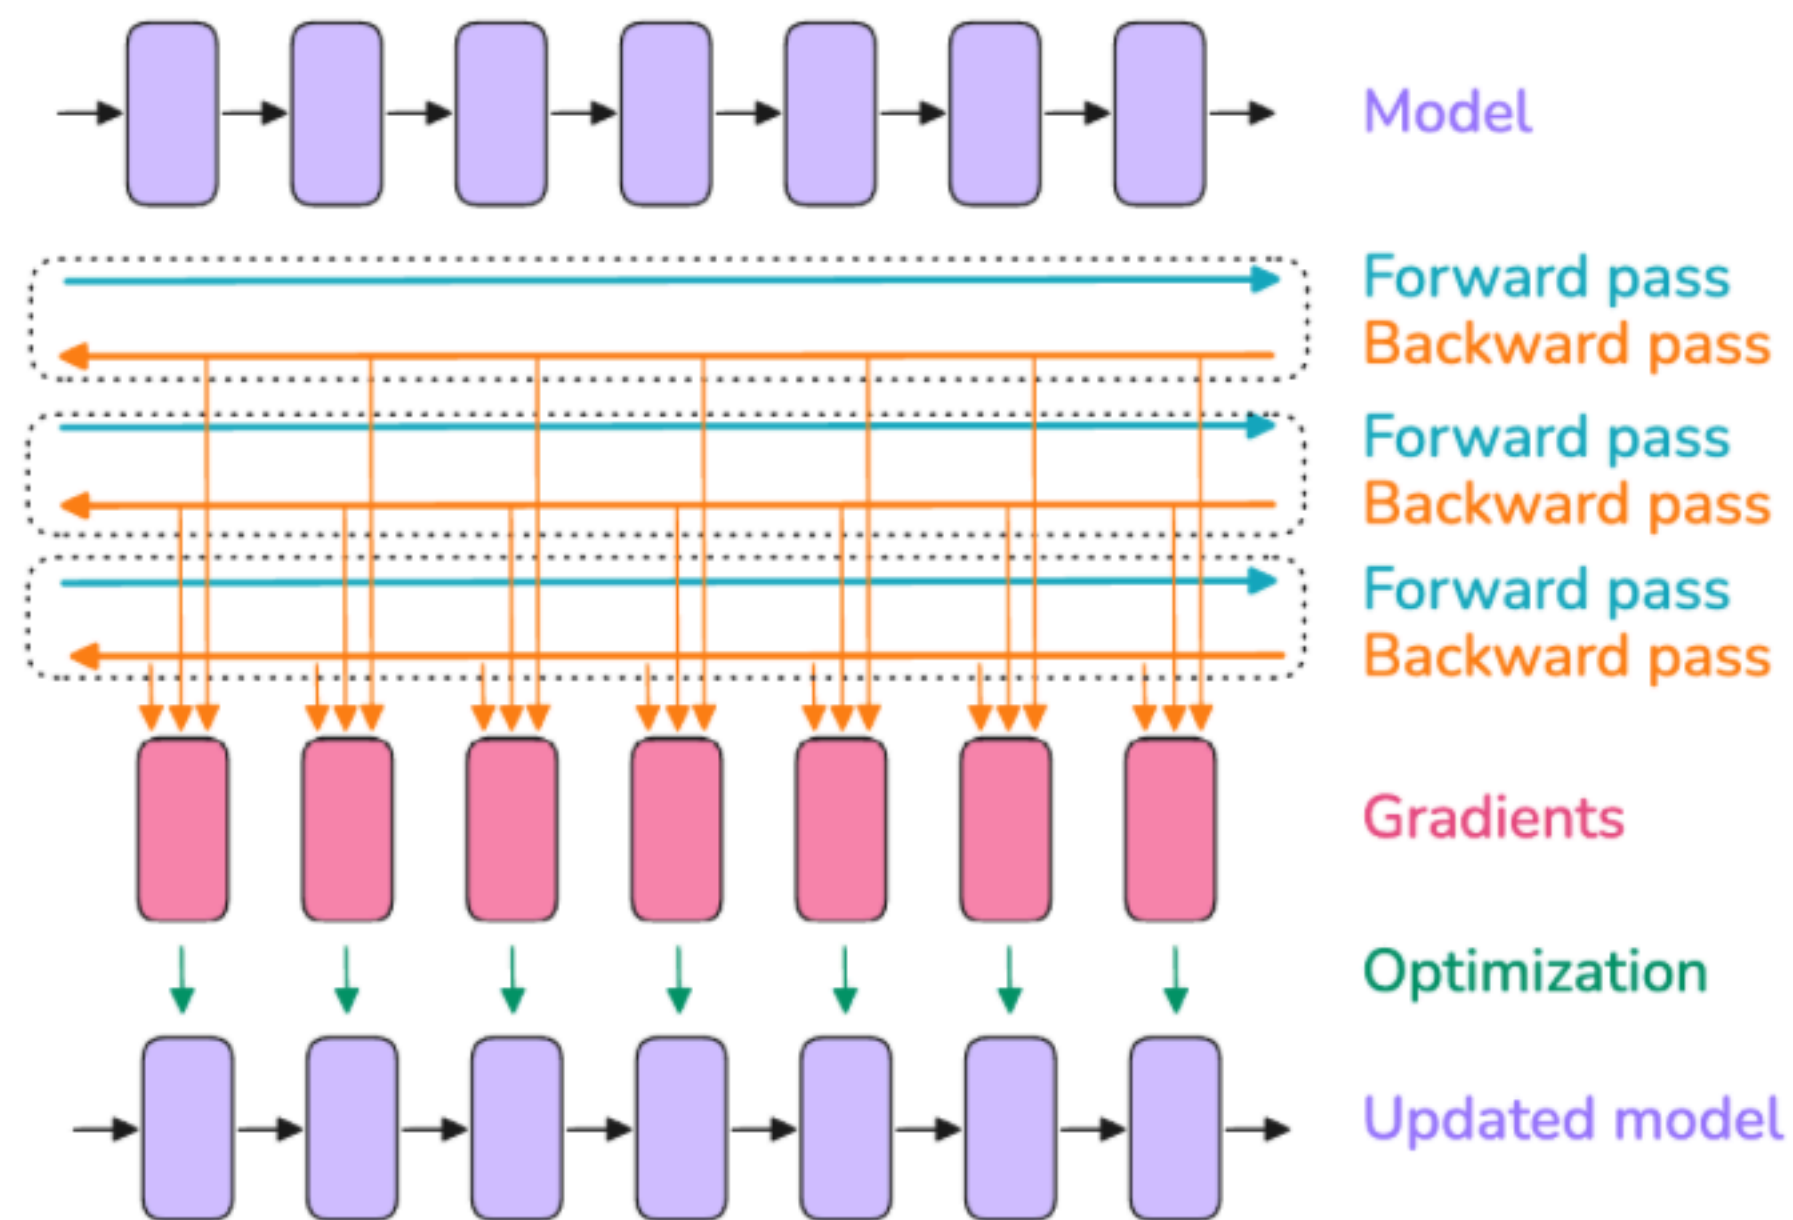


1. Run multiple forward-backward passes
2. Keeping a running mean of the gradients
3. When you hit the target bs then do optimizer step



# Gradient Accumulation

*Problem: we wish to run a batch size of 1000 but can only fit one sample in memory*



- *Run 1000 times:*
  - *Run forward-backward pass on 1 sample*
  - *Accumulate gradient in buffer, ie,  $grad += current\_grad / accumulation\_steps$*

We spent more compute (ran 1000 fwd-bwd passes) before doing an optimizer step to save memory

Key concept: Trade off memory for compute



# Recap

- Calculating FLOPs in transformer layers for inference vs training
- Understanding compute and memory in LLM training workloads
- Fitting things on one GPU: trading off memory for compute
  - Activation checkpointing
  - Gradient accumulation

# Why Is Scaling Hard?

```
`python train.py --small-model` -> `python train.py --big-model`
```

DUH!

Achieving “strong scaling”, ie, increase the number of chips used for training while achieving a proportional, linear increase in throughput, is hard due to communication overheads

Fitting everything (model, optimizer, gradients, activations) in memory

Choosing the right strategy for sharding / parallelizing when things don't fit in memory

Alleviating communication bottlenecks that arise from parallelisms and sharding

# Going Beyond A Single GPU

Let's assume we can model, gradients, optimizer, activation in one chip's memory

However, we have many chips lying around. How can we leverage these to get higher throughput?

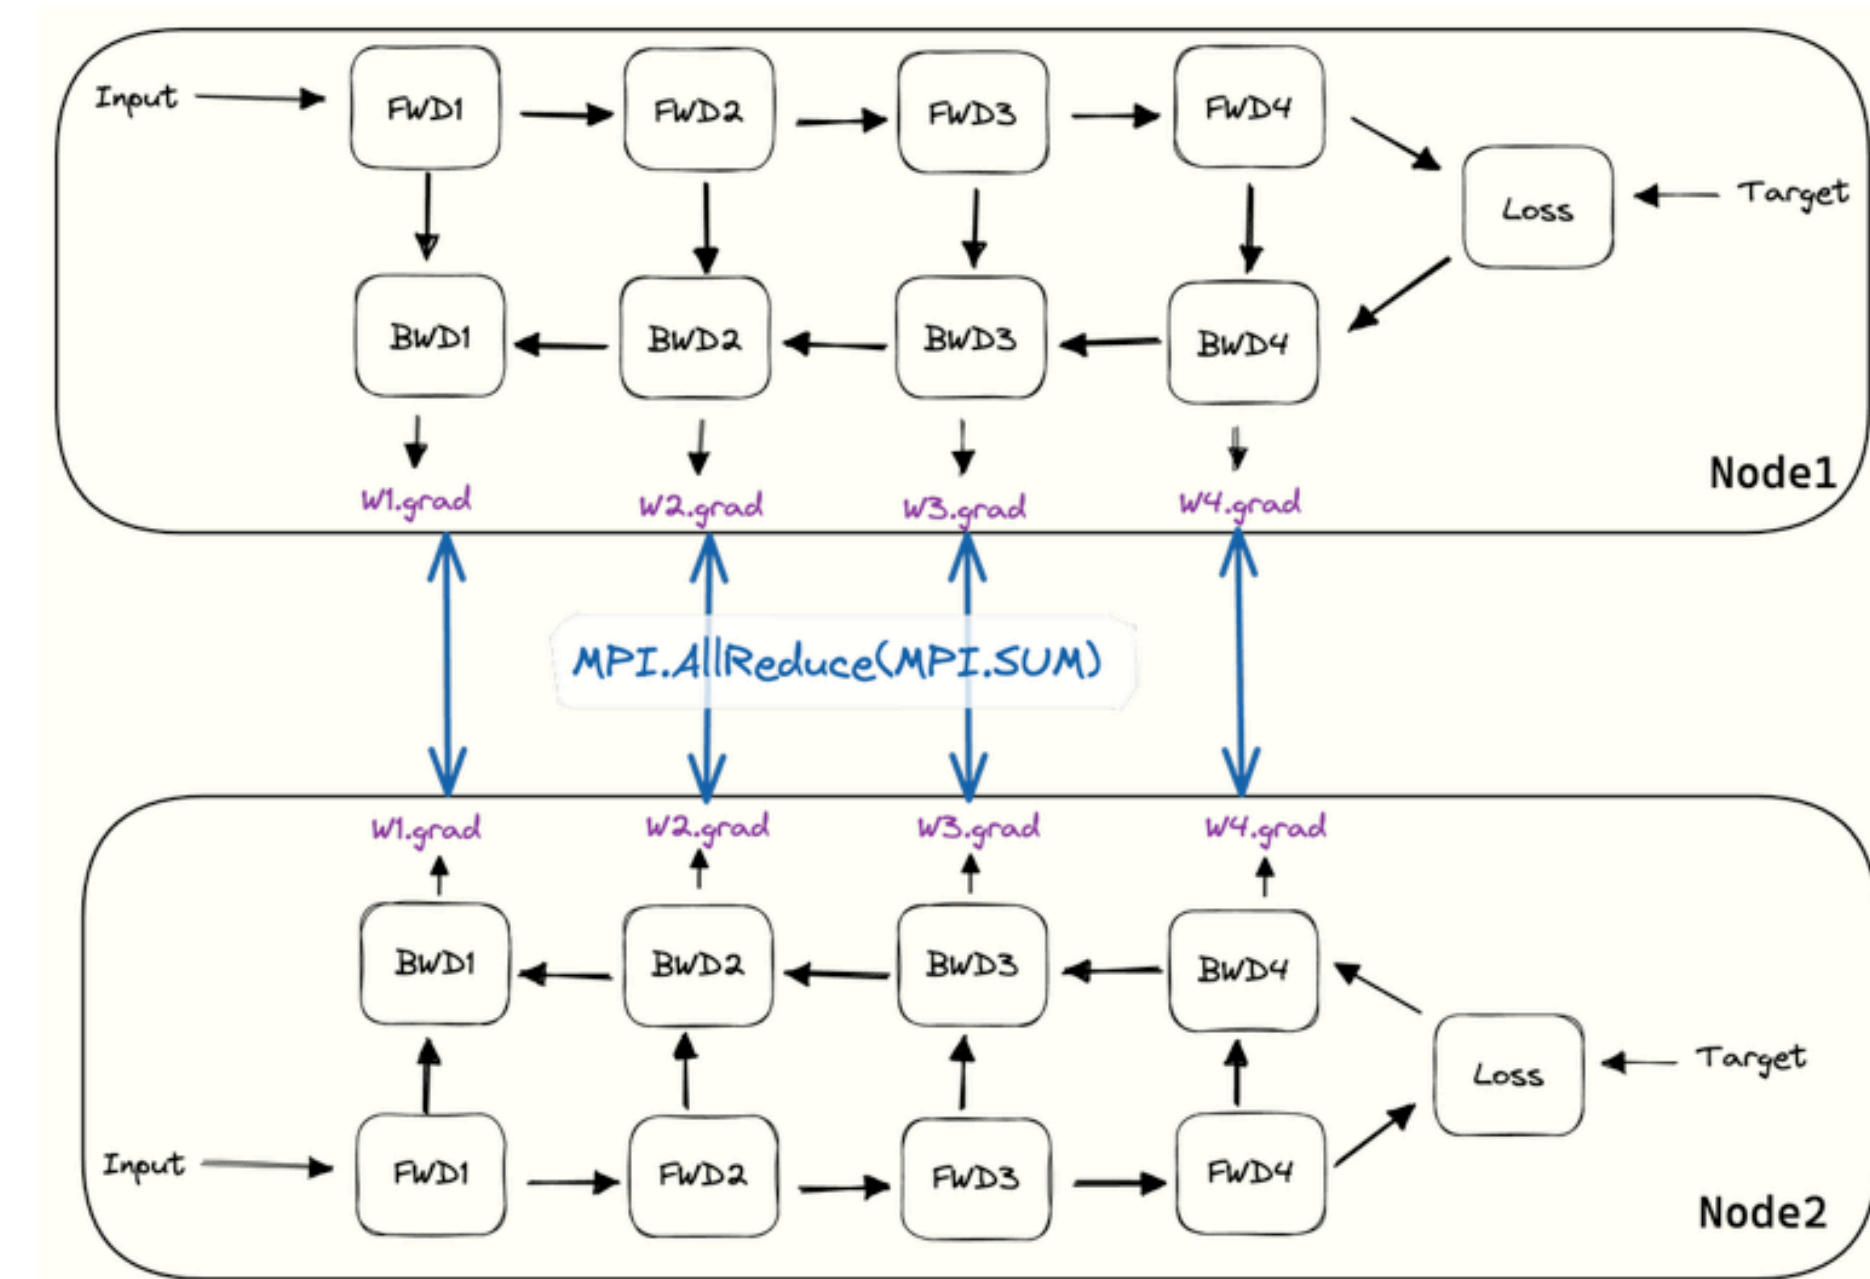
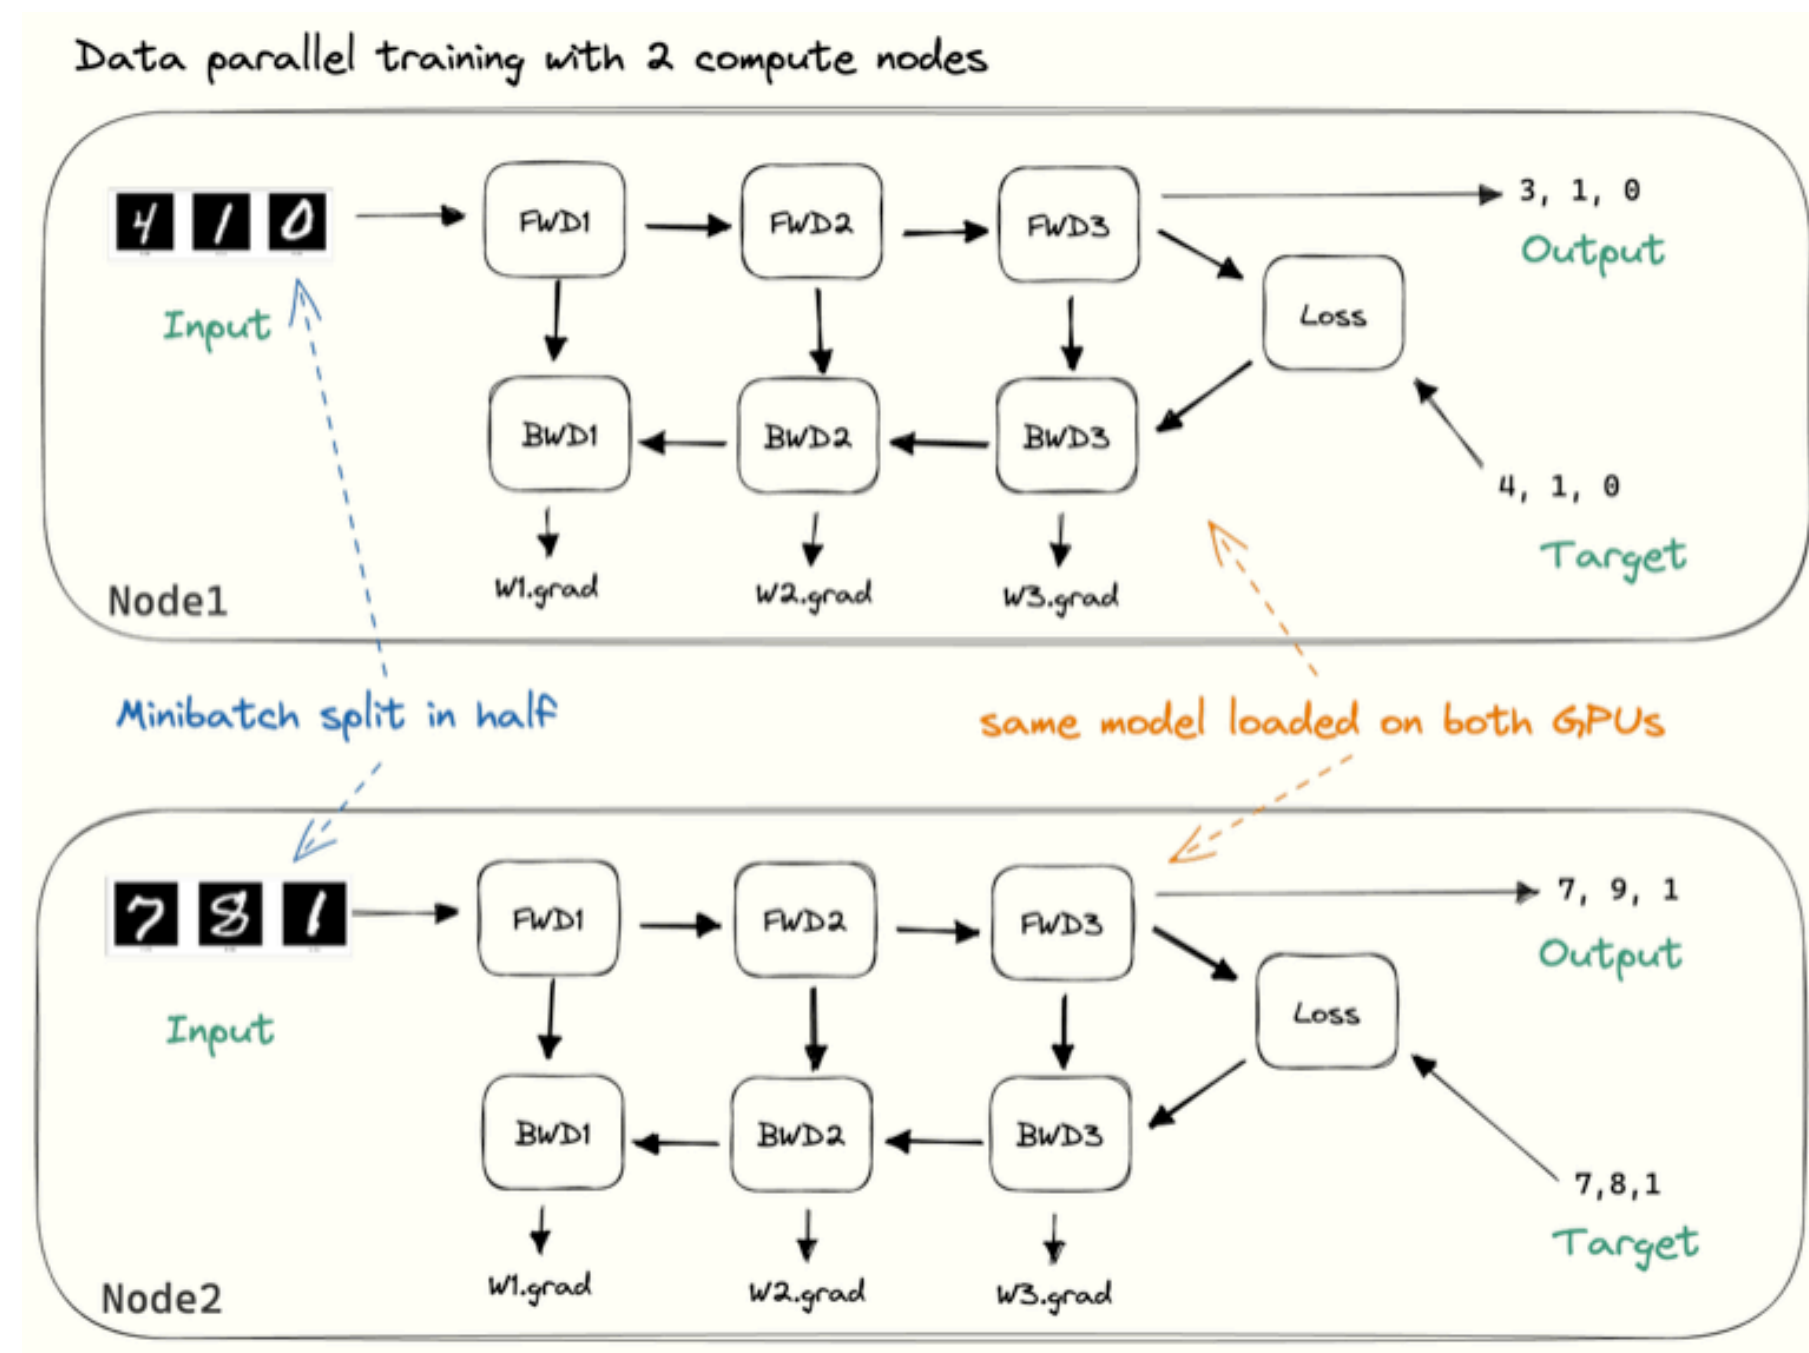
**Data parallelism: run different batches of data in parallel on different chips**

**Corollary: For this to be in sync**

- 1. Model weights must be duplicated on different chips**
- 2. After gradients are computed on different slices of data, they must be communicated across different GPUs**

# Data Parallelism

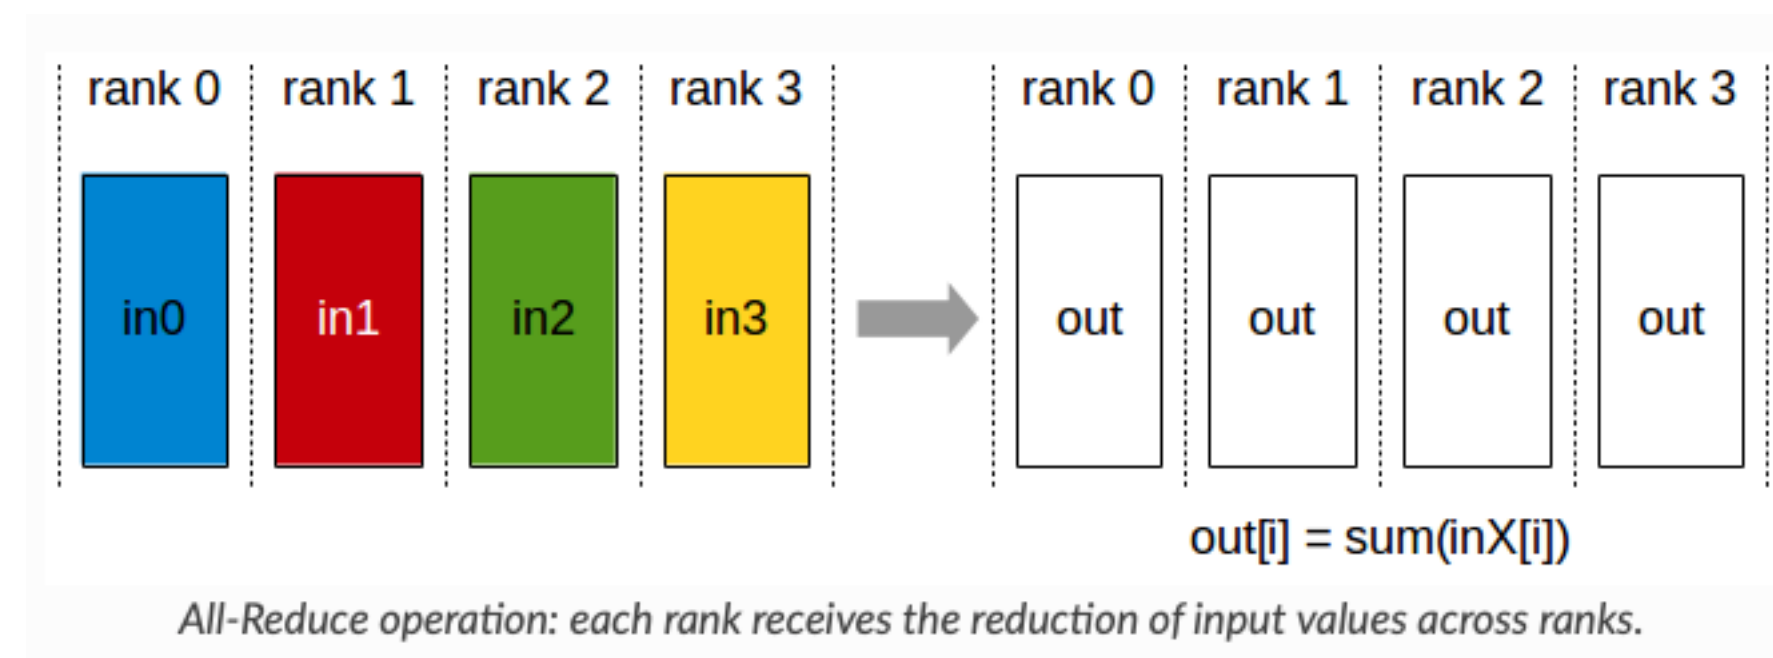
**Data parallelism: run different batches of data in parallel on different chips**



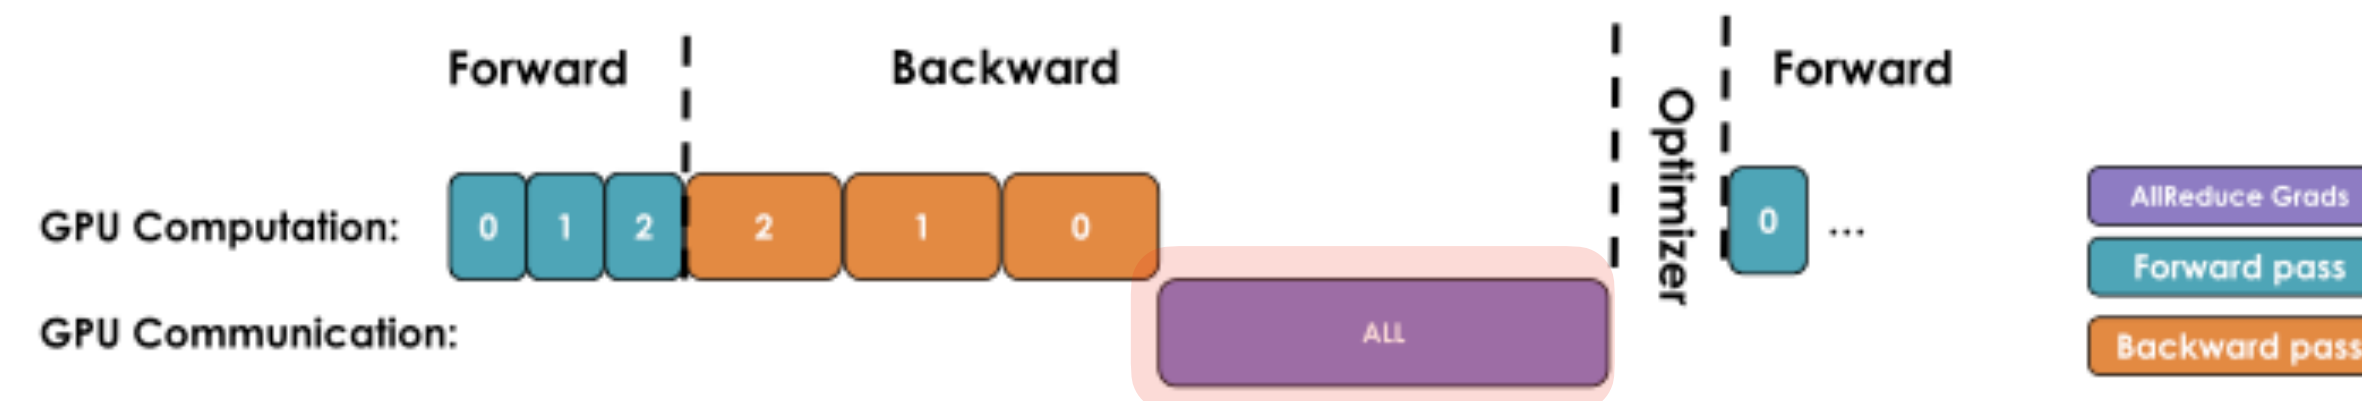
# Data Parallelism: Communication Overhead

**Data parallelism: run different batches of data in parallel on different chips**

Requires an all-reduce\* after gradients have been computed on each chip



View on each chip:

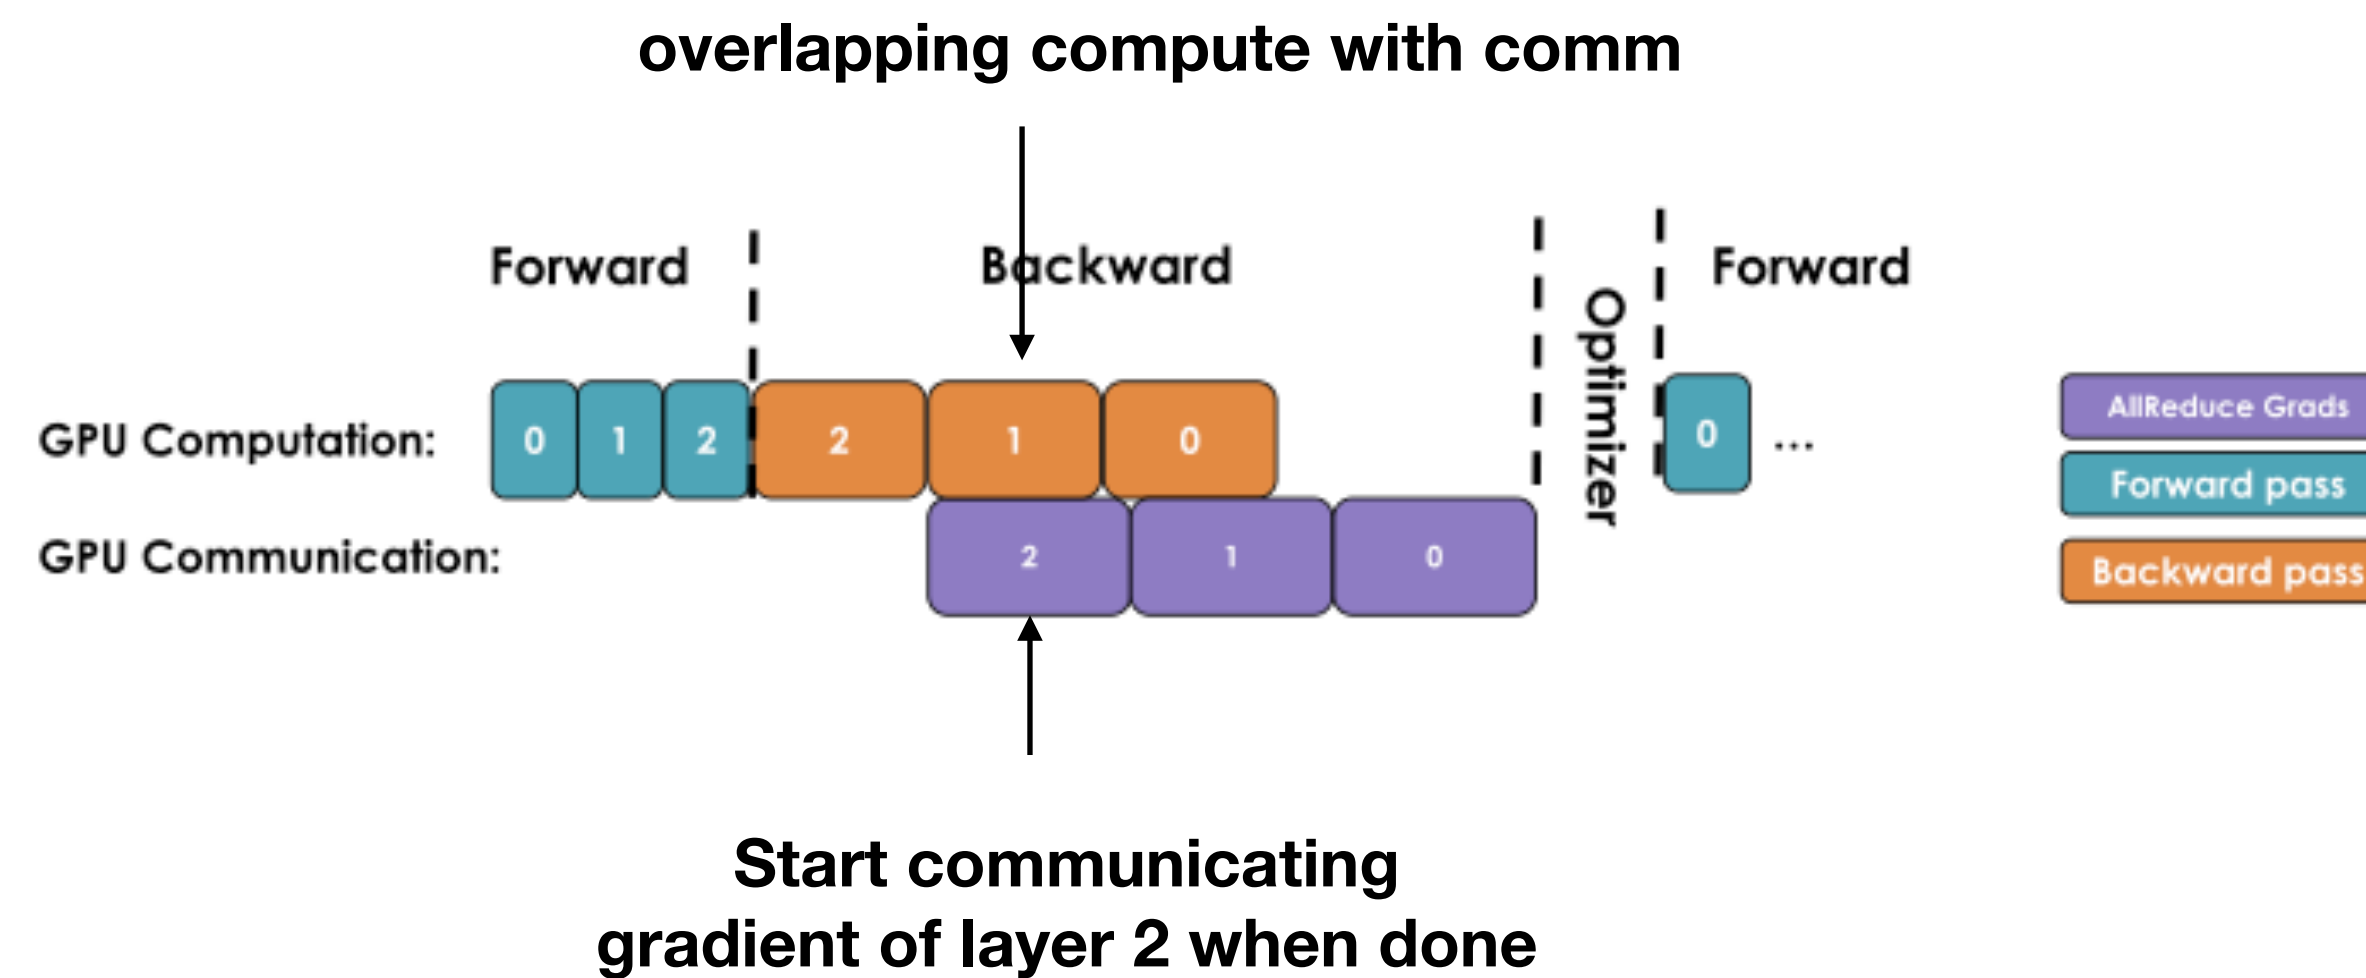


**Every device is waiting for communication**

\* <https://docs.nvidia.com/deeplearning/nccl/user-guide/docs/usage/collectives.html>

# Data Parallelism: Can We Do Better?

**Key Idea: Overlap comms with compute, hide the bottleneck**



**More in Next Lecture!**

# Recap

- Calculating FLOPs in transformer layers for inference vs training
- Understanding compute and memory in LLM training workloads
- Fitting things on one GPU: trading off memory for compute
  - Activation checkpointing
  - Gradient accumulation
- Data parallelism: communication bottlenecks
  - Overlapping comms with compute to remove bottlenecks



# References

1. *Ultrascale Playbook, HuggingFace, <https://huggingface.co/spaces/nanotron/ultrascale-playbook>*
2. *How to Scale Your Model, <https://jax-ml.github.io/scaling-book>*
3. *<https://medium.com/tensorflow/fitting-larger-networks-into-memory-583e3c758ff9>*
4. *Training Deep Nets with Sublinear Memory Cost, <https://arxiv.org/abs/1604.06174>*
5. *Reducing Activation Recomputation in Large Transformer Models, <https://arxiv.org/abs/2205.05198>*
6. *<https://github.com/karpathy/nano-llama31>*